

CSE214 – Analysis of Algorithms

PhD Furkan Gözükar, Toros University

https://github.com/FurkanGozukara/CSE214_2018

Lecture 9

Greedy Algorithms

*Based on Cevdet David Matuszek's
Lecture Notes - Director of MCIT*



Optimization problems

- An **optimization problem** is one in which you want to find, not just *a* solution, but the *best* solution
- A “greedy algorithm” sometimes works well for optimization problems
- A **greedy algorithm** works in phases. At each phase:
 - You take the best you can get right now, without regard for future consequences
 - You hope that by choosing a *local* optimum at each step, you will end up at a *global* optimum



Example: Counting money

- Suppose you want to count out a certain amount of money, using the fewest possible bills and coins
- A greedy algorithm would do this would be:
At each step, take the largest possible bill or coin that does not overshoot
 - Example: To make \$6.39, you can choose:
 - a \$5 bill
 - a \$1 bill, to make \$6
 - a 25¢ coin, to make \$6.25
 - A 10¢ coin, to make \$6.35
 - four 1¢ coins, to make \$6.39
- For US money, the greedy algorithm always gives the optimum solution



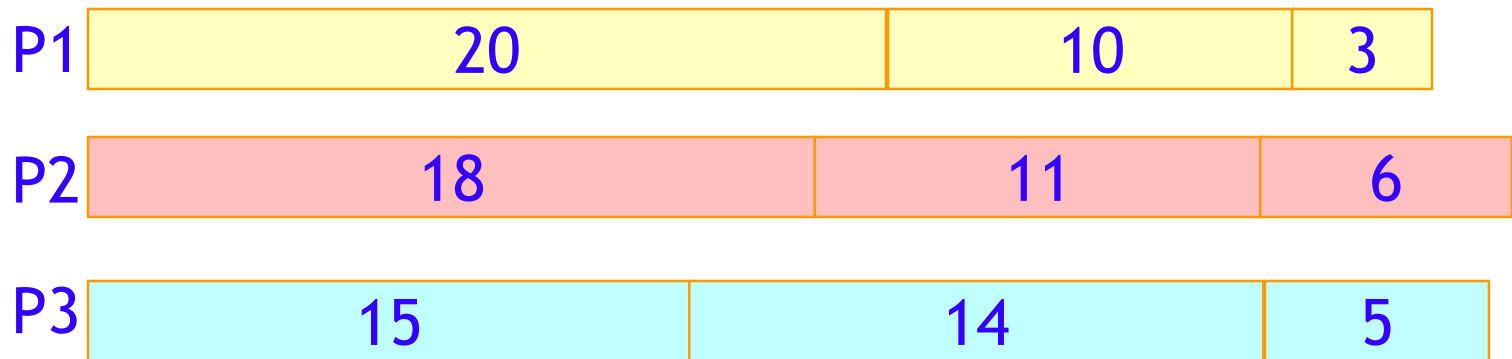
A failure of the greedy algorithm

- In some (fictional) monetary system, “krons” come in 1 kron, 7 kron, and 10 kron coins
- Using a greedy algorithm to count out 15 krons, you would get
 - A 10 kron piece
 - Five 1 kron pieces, for a total of 15 krons
 - This requires six coins
- A better solution would be to use two 7 kron pieces and one 1 kron piece
 - This only requires three coins
- The greedy algorithm results in a solution, but not in an optimal solution



A scheduling problem

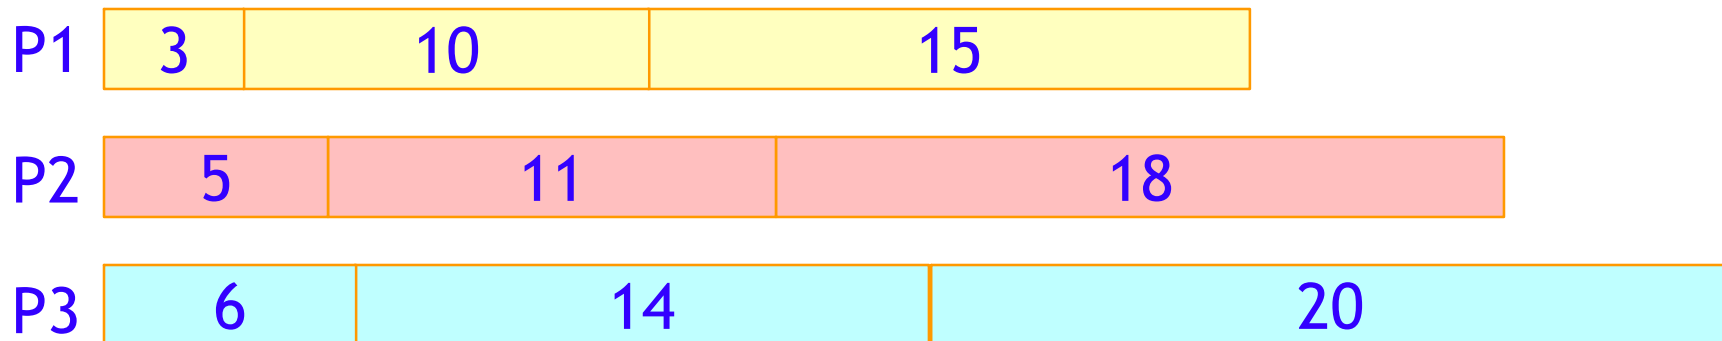
- You have to run nine jobs, with running times of 3, 5, 6, 10, 11, 14, 15, 18, and 20 minutes
- You have three processors on which you can run these jobs
- You decide to do the longest-running jobs first, on whatever processor is available



- Time to completion: $18 + 11 + 6 = 35$ minutes
- This solution isn't bad, but we might be able to do better

Another approach

- What would be the result if you ran the *shortest* job first?
- Again, the running times are 3, 5, 6, 10, 11, 14, 15, 18, and 20 minutes

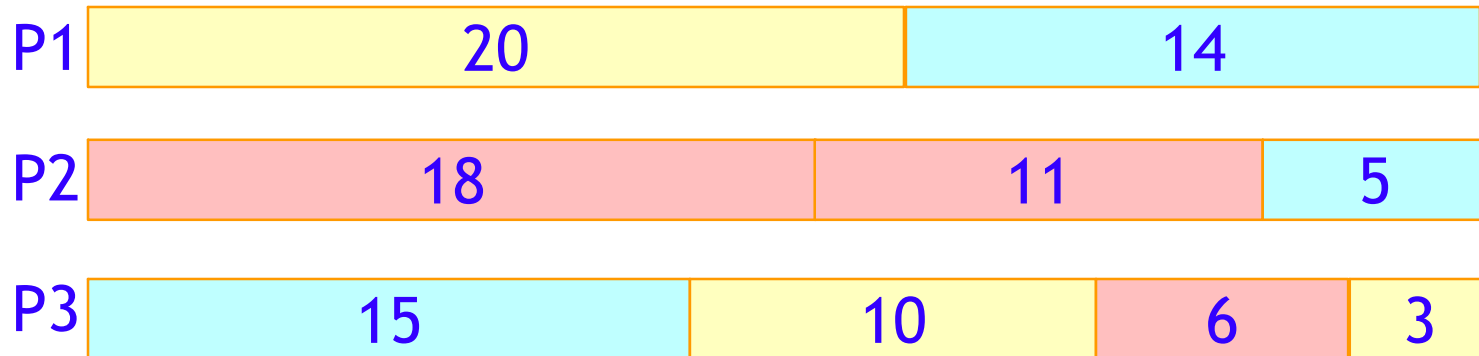


- That wasn't such a good idea; time to completion is now $6 + 14 + 20 = 40$ minutes
- Note, however, that the greedy algorithm itself is fast
 - All we had to do at each stage was pick the minimum or maximum



An optimum solution

- Better solutions do exist:



- This solution is clearly optimal (why?)
- Clearly, there are other optimal solutions (why?)
- How do we find such a solution?
 - One way: Try all possible assignments of jobs to processors
 - Unfortunately, this approach can take exponential time



Huffman Coding

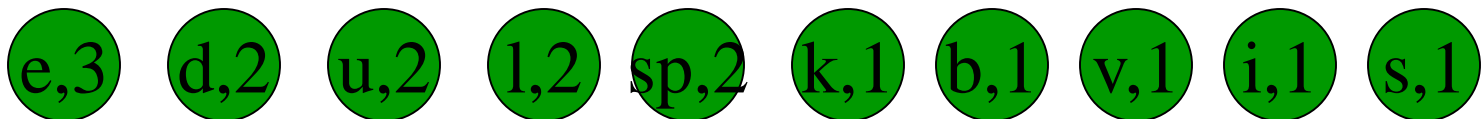
Prepared by Dan Stevenson
University of Wisconsin

- Huffman codes can be used to compress information
 - Like WinZip – although WinZip doesn't use the Huffman algorithm
 - JPEGs do use Huffman as part of their compression process
- The basic idea is that instead of storing each character in a file as an 8-bit ASCII value, we will instead store the more frequently occurring characters using fewer bits and less frequently occurring characters using more bits
 - On average this should decrease the filesize (usually $\frac{1}{2}$)



Huffman Coding

- As an example, lets take the string:
 "duke blue devils"
- We first to a frequency count of the characters:
 - e:3, d:2, u:2, l:2, space:2, k:1, b:1, v:1, i:1, s:1
- Next we use a Greedy algorithm to build up a Huffman Tree
 - We start with nodes for each character





Huffman Coding

- We then pick the nodes with the smallest frequency and combine them together to form a new node
 - The selection of these nodes is the Greedy part
- The two selected nodes are removed from the set, but replace by the combined node
- This continues until we have only 1 node left in the set

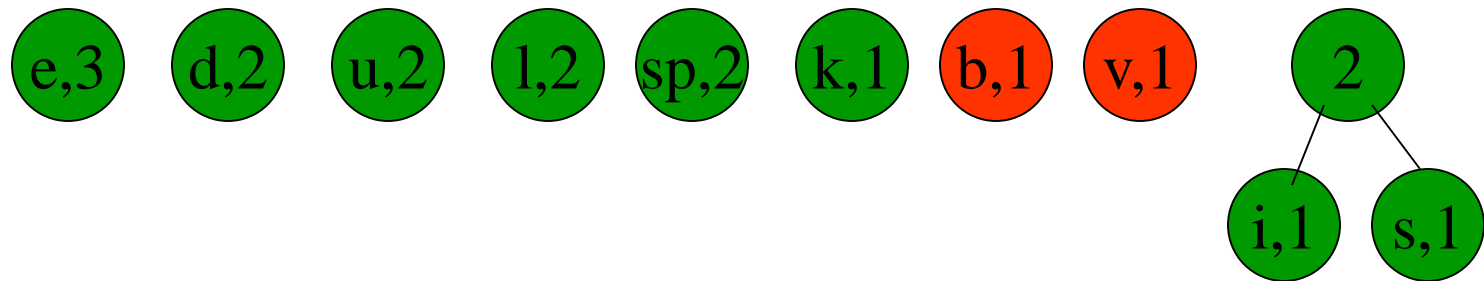


Huffman Coding

e,3 d,2 u,2 l,2 sp,2 k,1 b,1 v,1 i,1 s,1

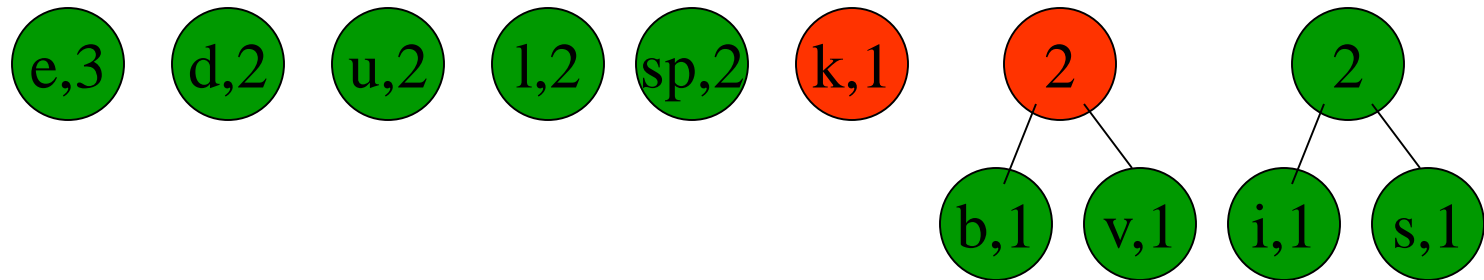


Huffman Coding



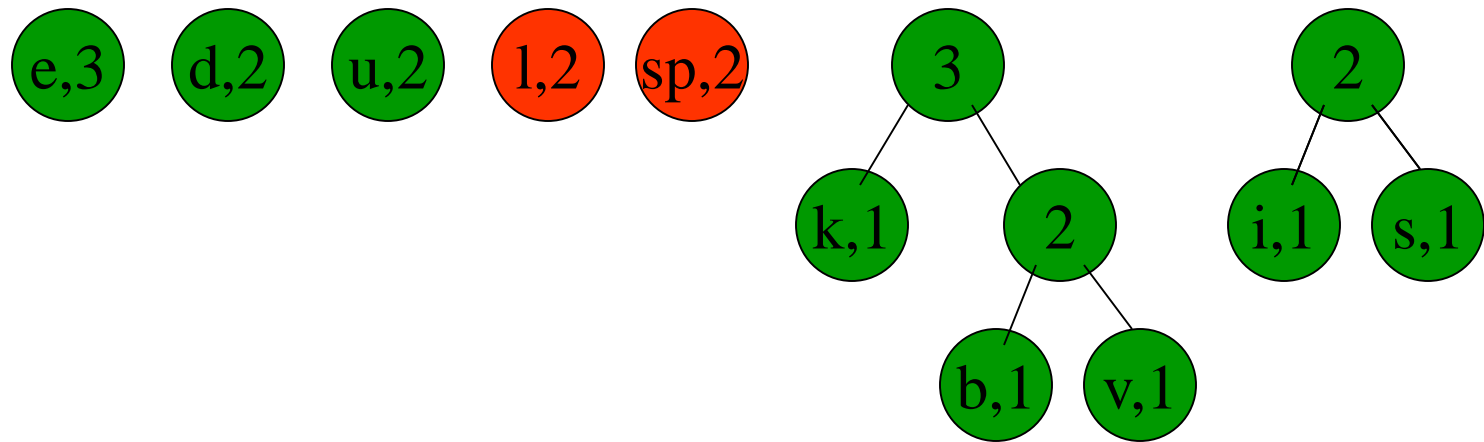


Huffman Coding

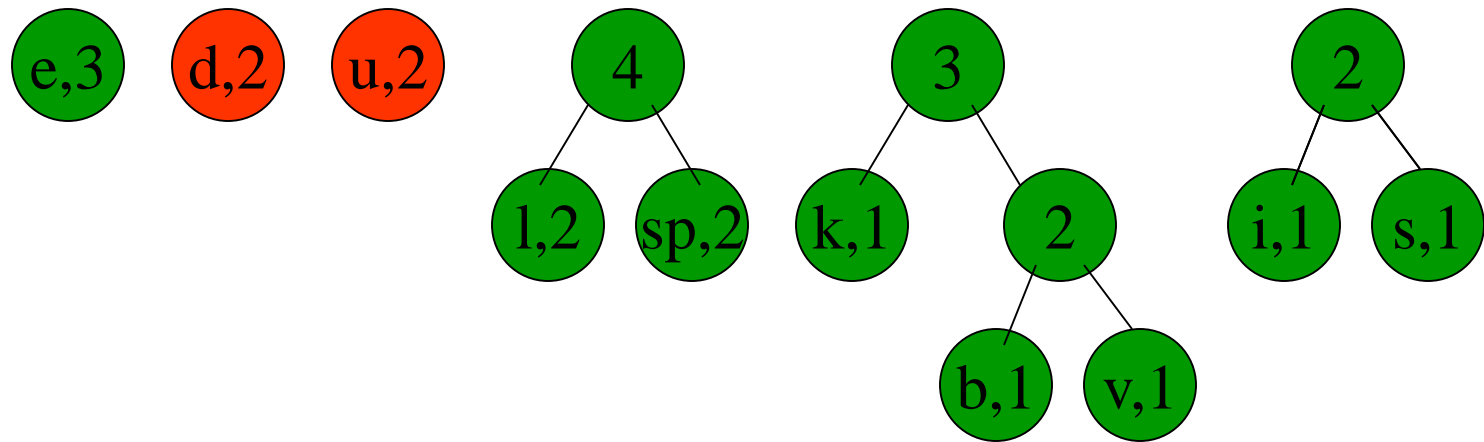




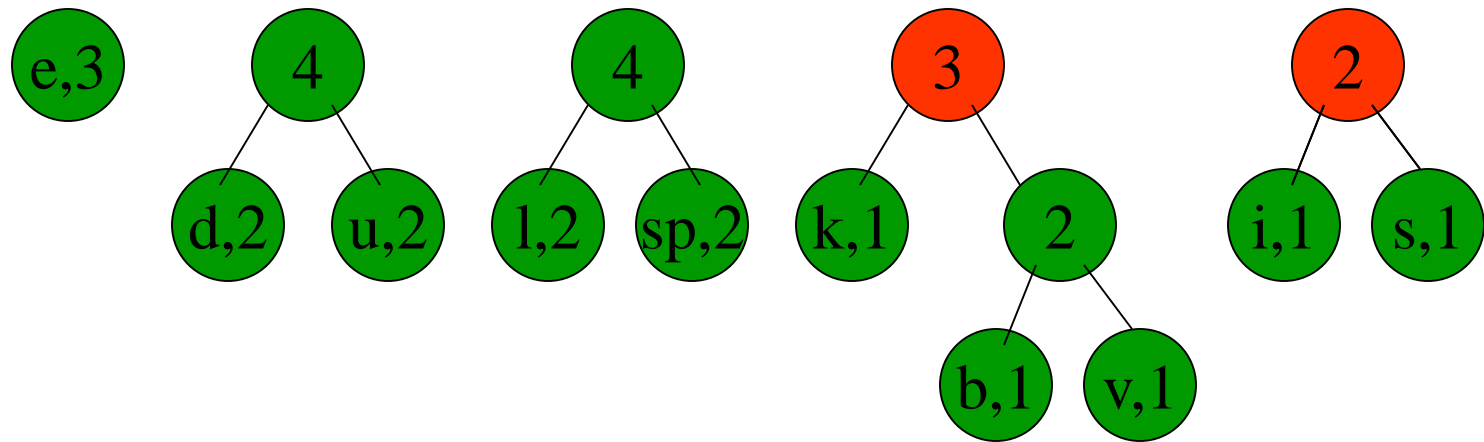
Huffman Coding



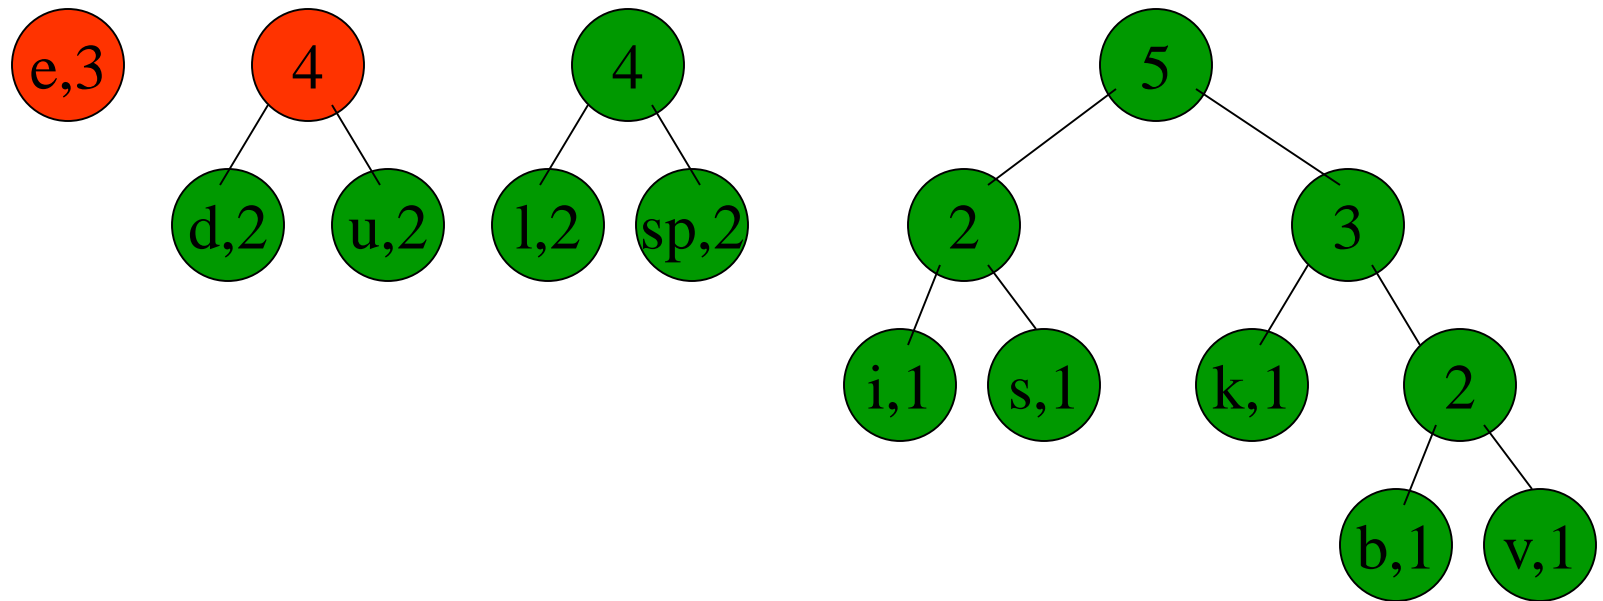
Huffman Coding



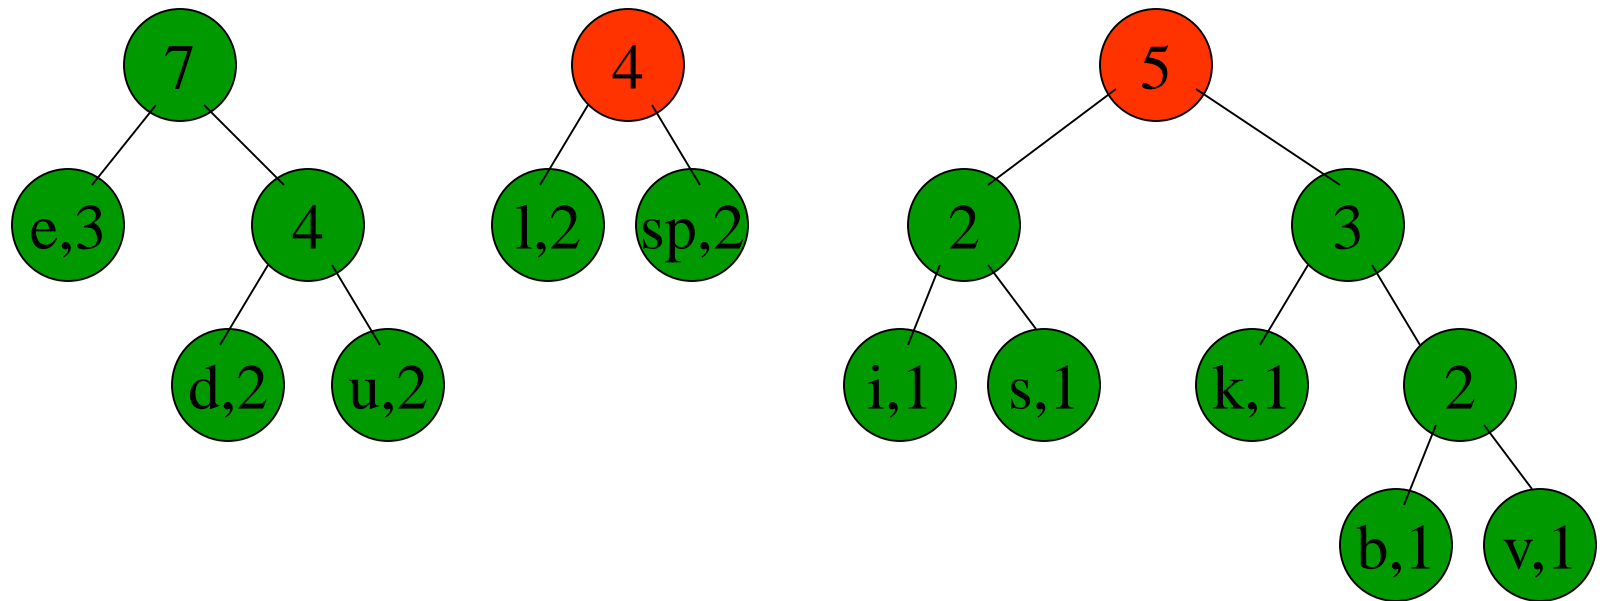
Huffman Coding



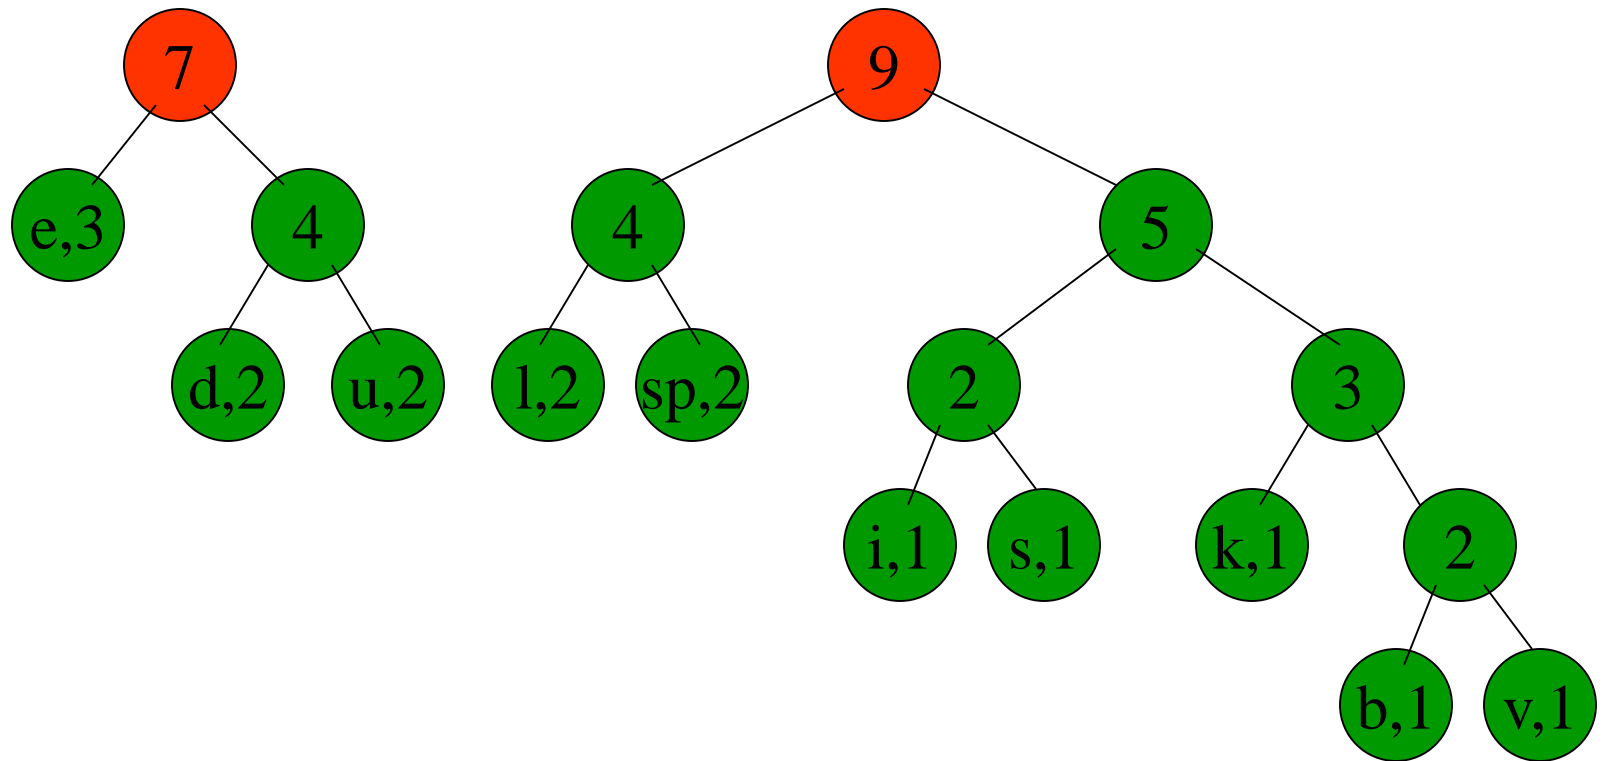
Huffman Coding



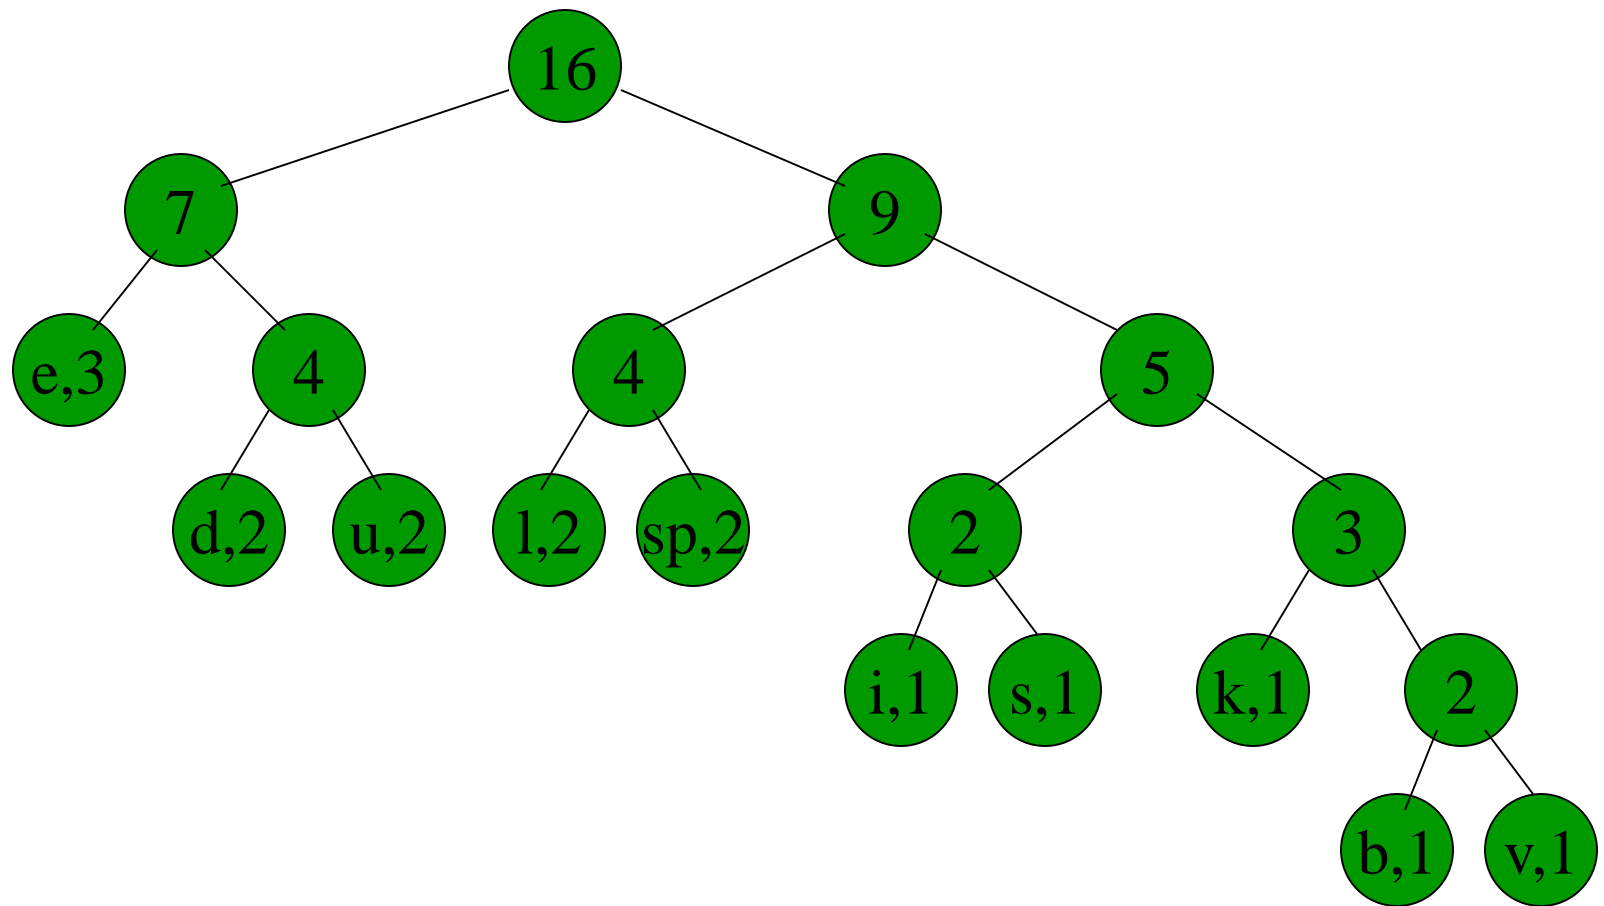
Huffman Coding



Huffman Coding



Huffman Coding

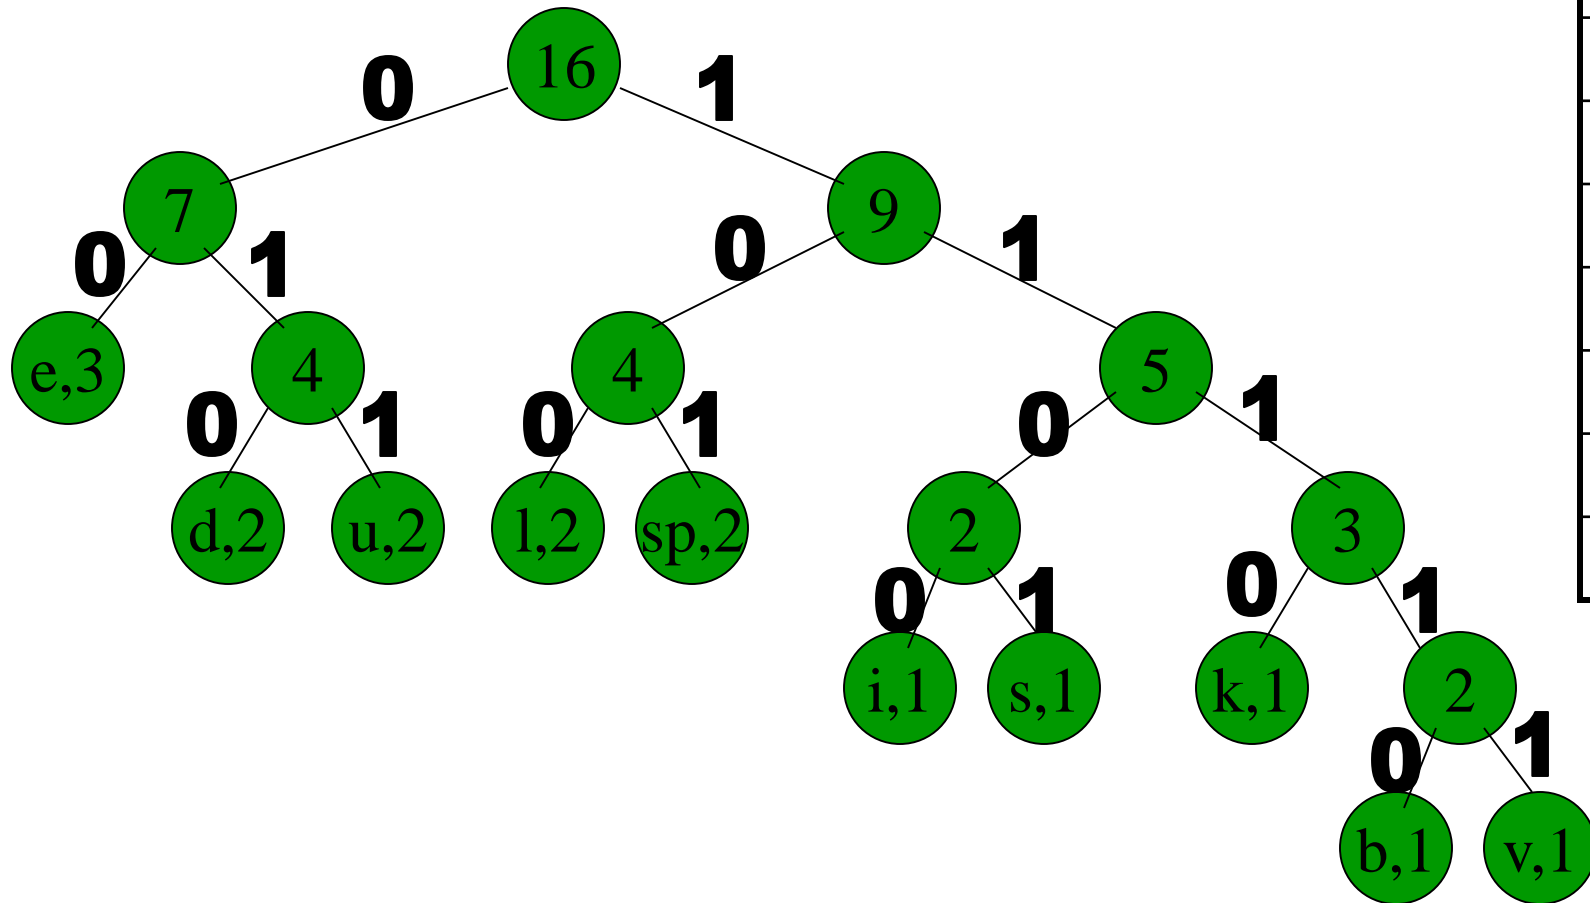




Huffman Coding

- Now we assign codes to the tree by placing a 0 on every left branch and a 1 on every right branch
- A traversal of the tree from root to leaf give the Huffman code for that particular leaf character
- Note that no code is the prefix of another code

Huffman Coding



e	00
d	010
u	011
l	100
sp	101
i	1100
s	1101
k	1110
b	11110
v	11111



Huffman Coding

- These codes are then used to encode the string

- Thus, “duke blue devils” turns into:

010 011 1110 00 101 11110 100 011 00 101 010 00 11111 1100 100 1101

- When grouped into 8-bit bytes:

01001111 10001011 11101000 11001010 10001111 11100100 1101xxxx

- Thus it takes 7 bytes of space compared to 16 characters * 1 byte/char = 16 bytes uncompressed

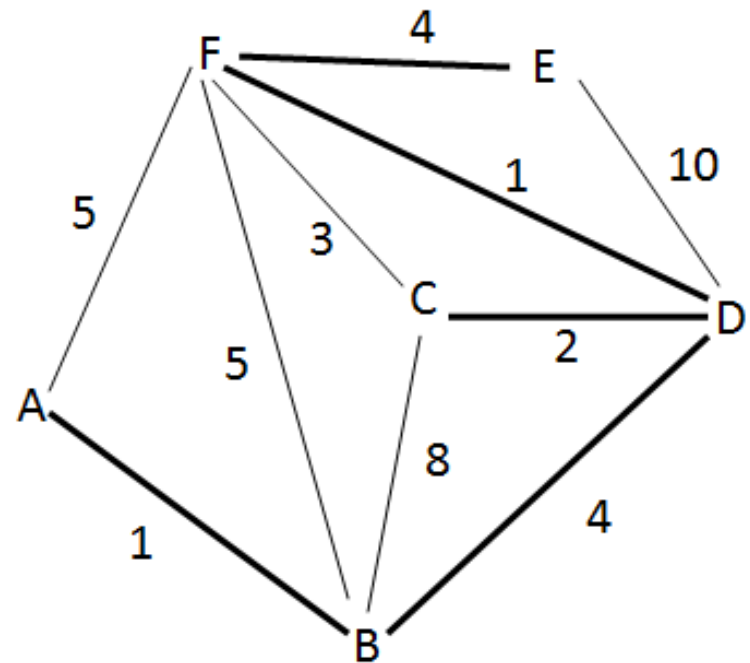


Huffman Coding

- Uncompressing works by reading in the file bit by bit
 - Start at the root of the tree
 - If a 0 is read, head left
 - If a 1 is read, head right
 - When a leaf is reached decode that character and start over again at the root of the tree
- Thus, we need to save Huffman table information as a header in the compressed file
 - Doesn't add a significant amount of size to the file for large files (which are the ones you want to compress anyway)
 - Or we could use a fixed universal set of codes/frequencies

Minimum Spanning Tree

- Input: graph G with weights on the edges
- Output: connected sub graph G' of G that includes all the vertices of G , of minimum total weight





Exhaustive Search

- List all connected sub-graphs of G
- Return sub-graph with least weight
- Number of vertices, N , and number of edges, M
- $O(m^{n-1})$

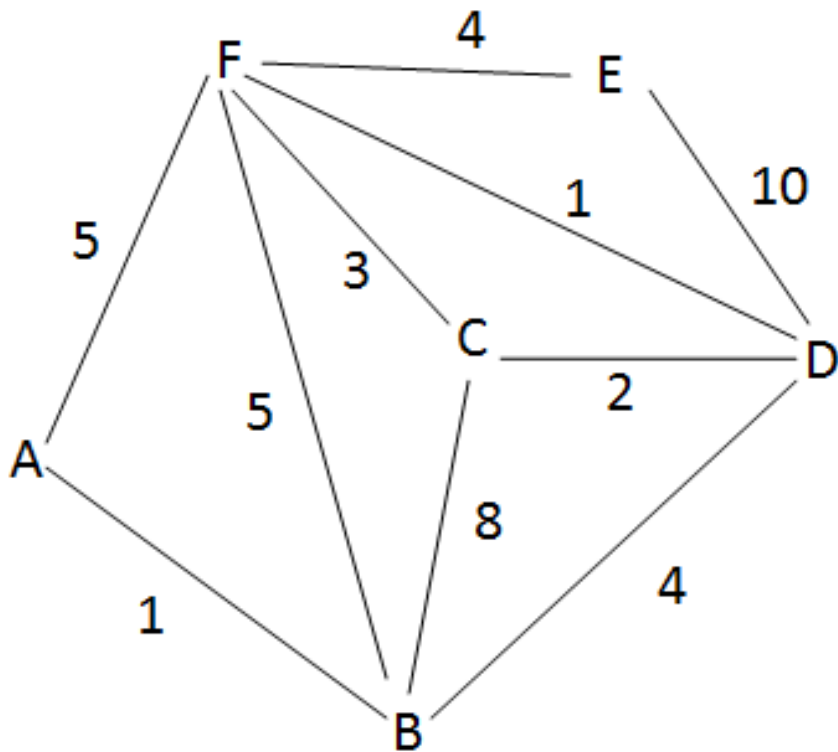


Greedy Algorithm

- Start with a graph containing just the vertices, $G' = \{V, \emptyset\}$
- Add edges with the least weight until the graph is connected but no cycles are created
- To do this:
 - Order the edges in increasing weight
 - While the graph is not connected, add an edge
 - End when all vertices are connected

Greedy Algorithm - Example

Initial Graph:



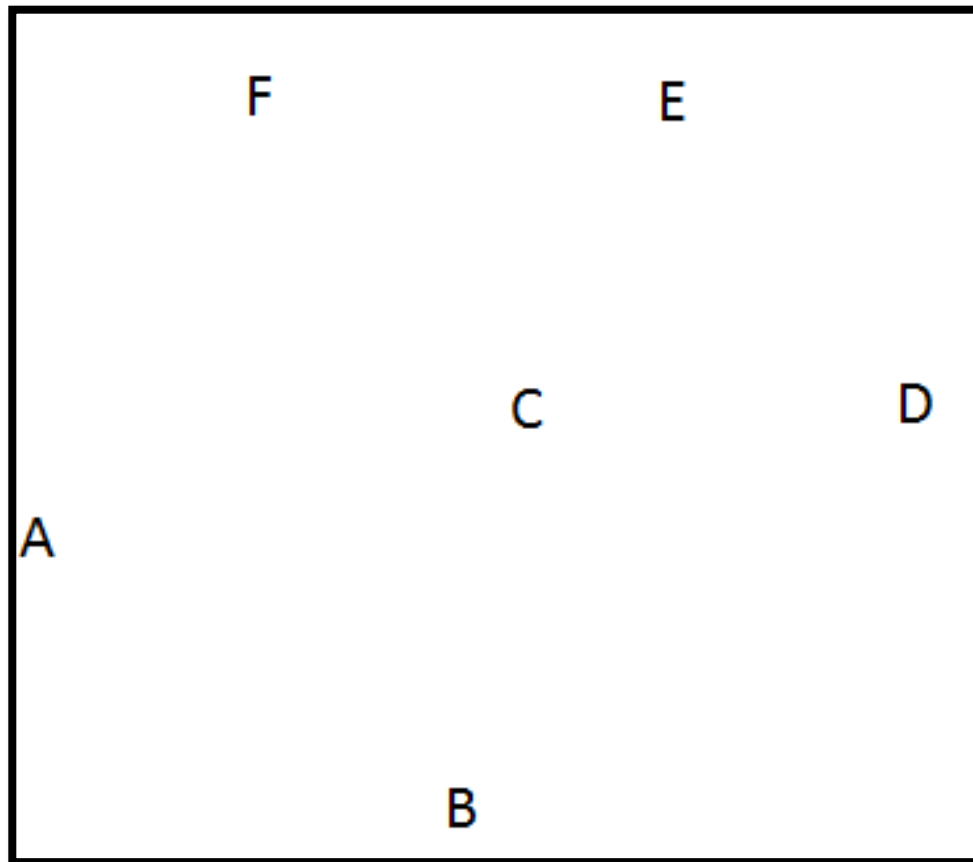
List of Edges:

Edge	Weight
AB	1
AF	5
BC	8
BD	4
BF	5
CD	2
CF	3
DE	10
DF	1
EF	4

Greedy Algorithm - Example

Start with $G' = \{V, \emptyset\}$ and sorted list of edges

Graph:



Sorted List of Edges:

Edge	Weight
AB	1
DF	1
CD	2
CF	3
BD	4
EF	4
AF	5
BF	5
BC	8
DE	10

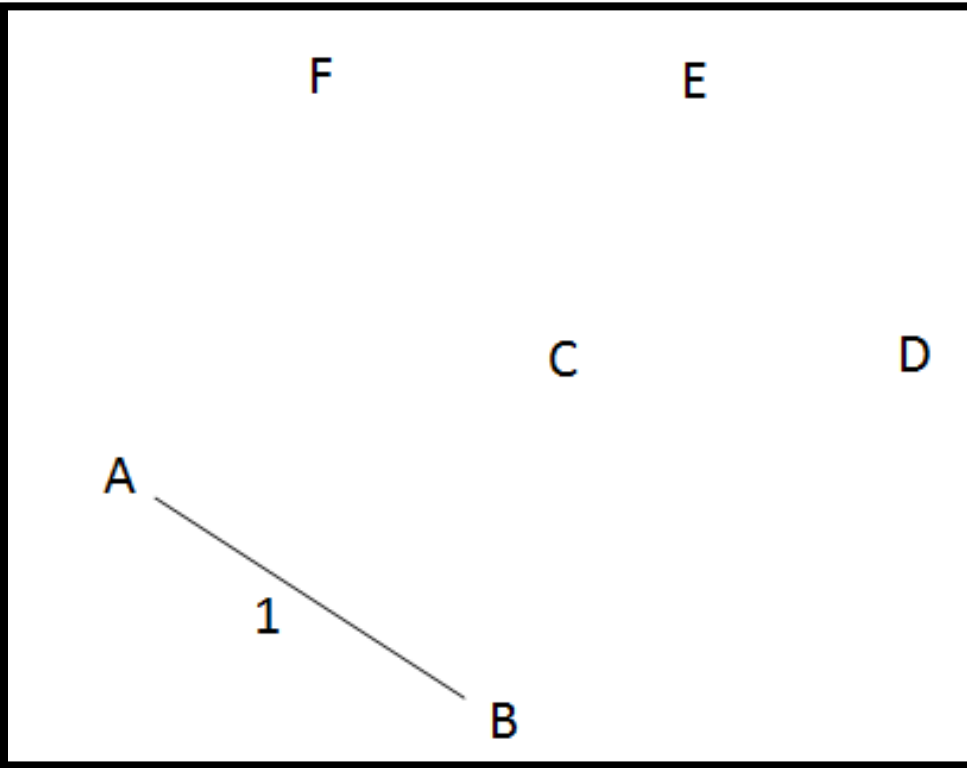
Greedy Algorithm - Example

Add edge from list to graph s.t. no cycles are created
Remove edge from list

Graph:

Sorted List of Edges:

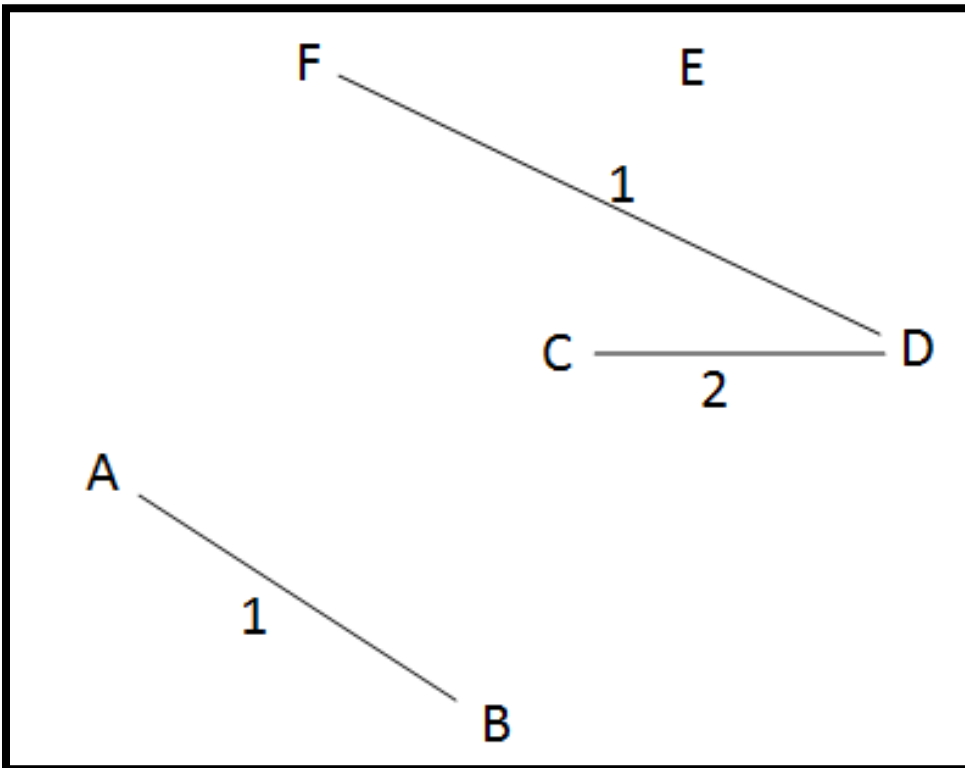
Edge	Weight
AB	4
DF	1
CD	2
CF	3
BD	4
EF	4
AF	5
BF	5
BC	8
DE	10



Greedy Algorithm - Example

Add edge from list to graph s.t. no cycles are created
Remove edge from list

Graph:



Sorted List of Edges:

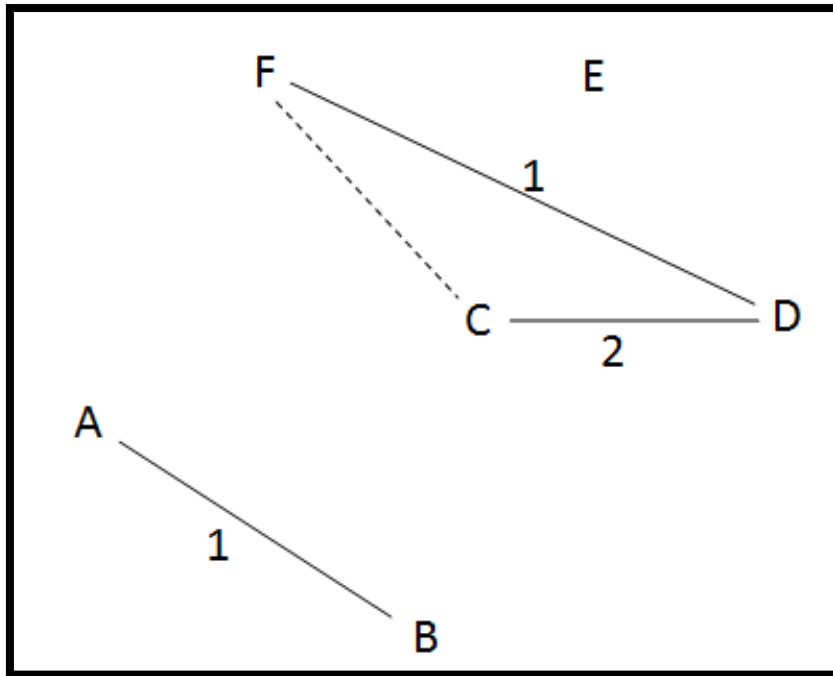
Edge	Weight
AB	1
DF	1
CD	2
CF	3
BD	4
EF	4
AF	5
BF	5
BC	8
DE	10

Greedy Algorithm - Example

Add edge from list to graph s.t. no cycles are created

Remove edge from list

Graph:



Sorted List of Edges:

Edge	Weight
CF	3
BD	4
EF	4
AF	5
BF	5
BC	8
DE	10

Check for cycles using Depth First Search starting at a vertex in the added edge:

- Start at C
- C → D → F → C
- C gets visited twice ⇒ a cycle exists and CF should not be added to the graph

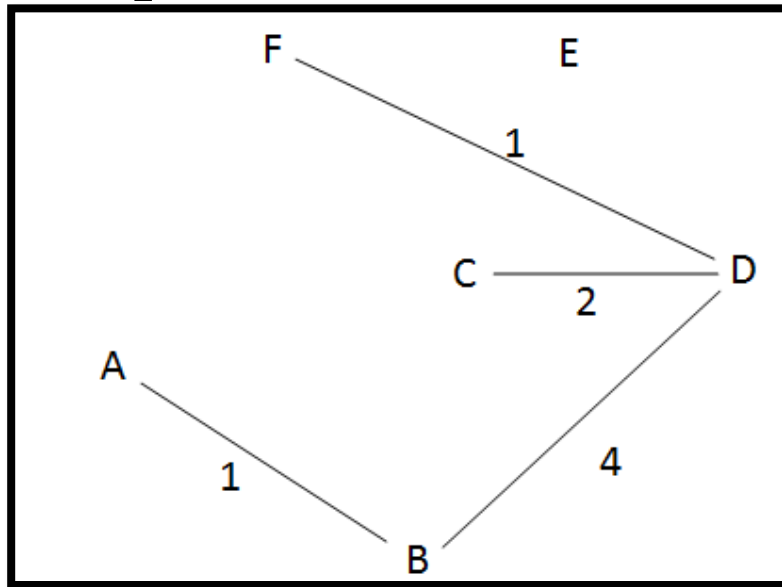
Greedy Algorithm - Example

Add edge from list to graph s.t. no cycles are created

Remove edge from list

End when all vertices can be visited (graph is connected)

Graph:



Sorted List of Edges:

Edge	Weight
BD	4
EF	4
AF	5
BF	5
BC	8
DE	10

Check for cycles using Depth First Search starting at a vertex in the added edge:

- Start at B
- B → D → C → F → A
- If all vertices are visited by DFS, then the graph is connected and we are done

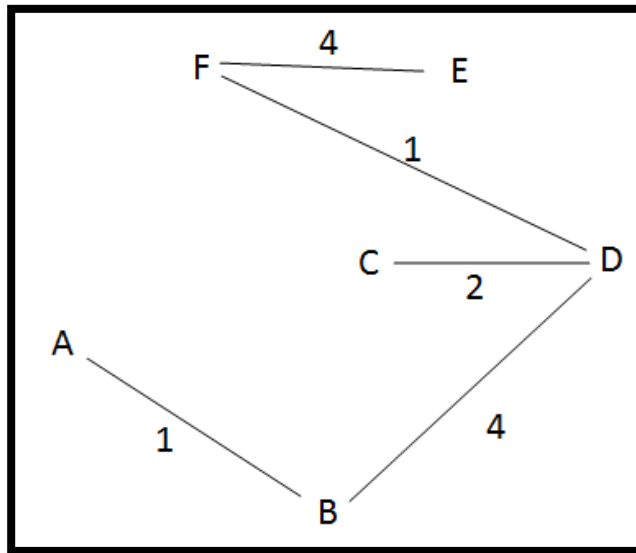
Greedy Algorithm - Example

Add edge from list to graph s.t. no cycles are created

Remove edge from list

End when all vertices can be visited (graph is connected)

Graph:



Sorted List of Edges:

Edge	Weight
EF	4
AF	5
BF	5
BC	8
DE	10

Check for cycles using Depth First Search starting at a vertex in the added edge:

- Start at E
- E -> F -> D -> B -> A -> C
- All vertices were visited and there were no cycles => we have found a minimum spanning tree with weight 12.



Greedy Algorithm- Pseudocode

- Given $G = (V, E)$
- $G' \leftarrow (V, \emptyset)$
- While G' is not connected
 - Add $e \in E$ to G' s.t. G' is acyclic and e is minimum
 - Remove e from E



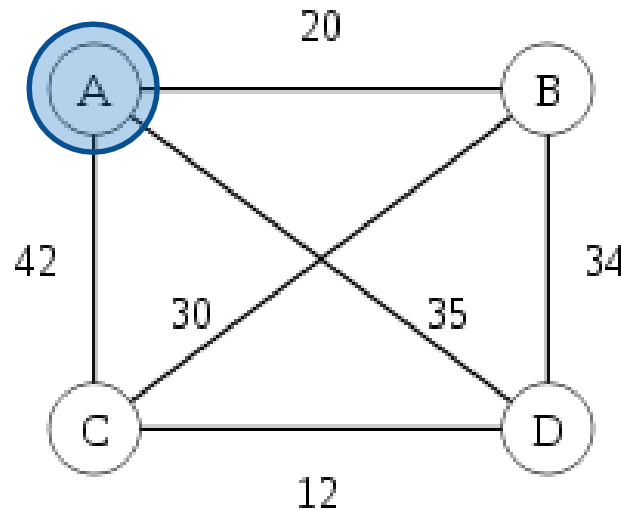
Greedy Algorithm – Run Time

- Initialization – constant
- While loop – $O(n-1)$
 - Connected graph has $n-1$ edges
 - Must add $n-1$ edges to the graph for it to be connected
- Find a minimum $e \in E$ – $O(m)$
- Make sure G' is acyclic – $O(2n)$
 - DFS is $O(m+n) < O(2n)$ where $m \leq n-1$
- Test connectivity of G' – $O(n)$
 - Can use DFS; could be done in same step as testing acyclicity
- Remove e from E - constant
- Total Runtime: $O(n \cdot (m + 2n + n)) \sim O(nm + n^2) \rightarrow$
POLYNOMIAL

Traveling Salesman Problem (TSP)

Prepared by Caisar Oentoro

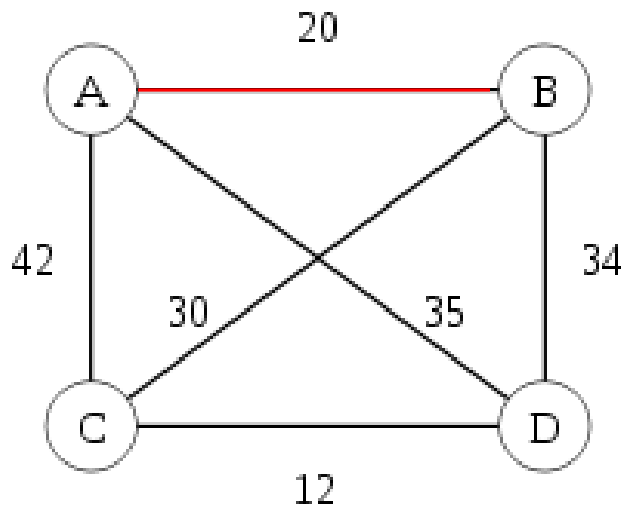
- > The Problem is how to travel from city A and visit all city on the map, then back to city A again.



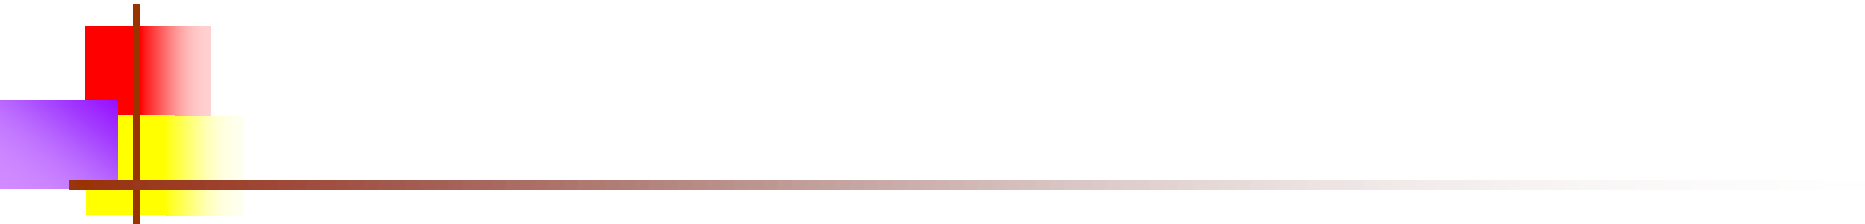
- > The rules: you only visit each city once and you can't pass through any traversed path.

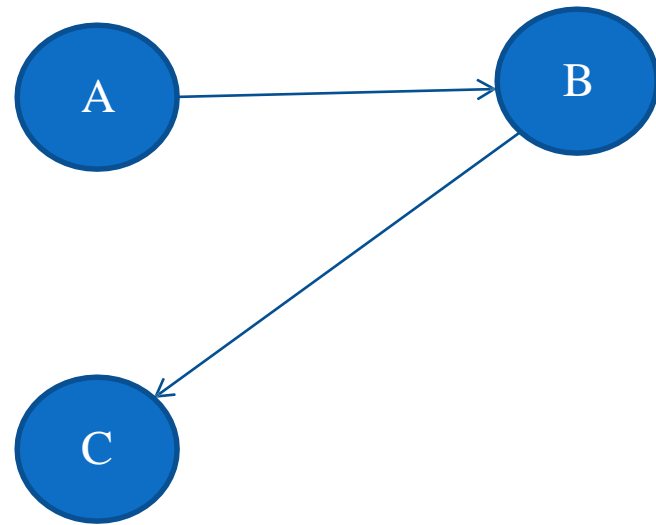
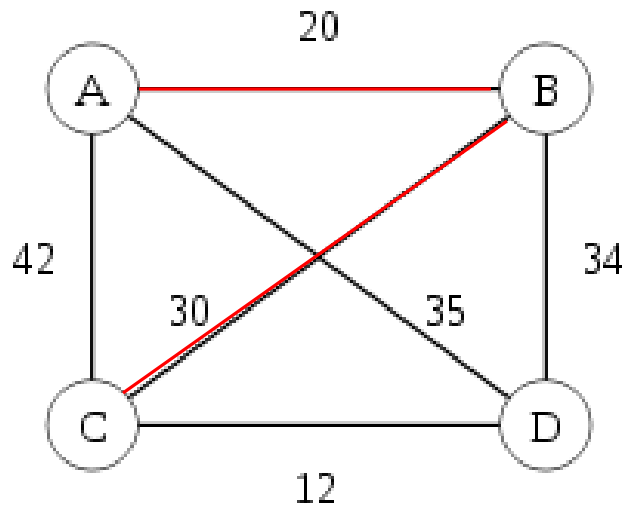
Solution:

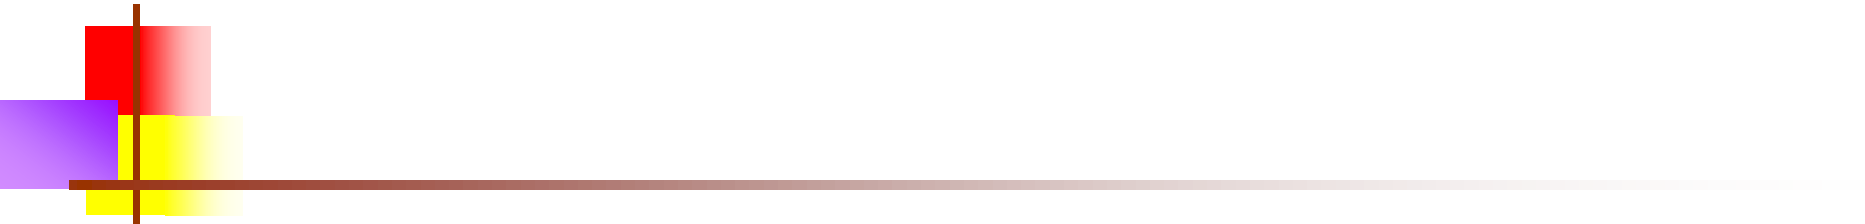
- > Find the shortest path from city A(start) to any other city.

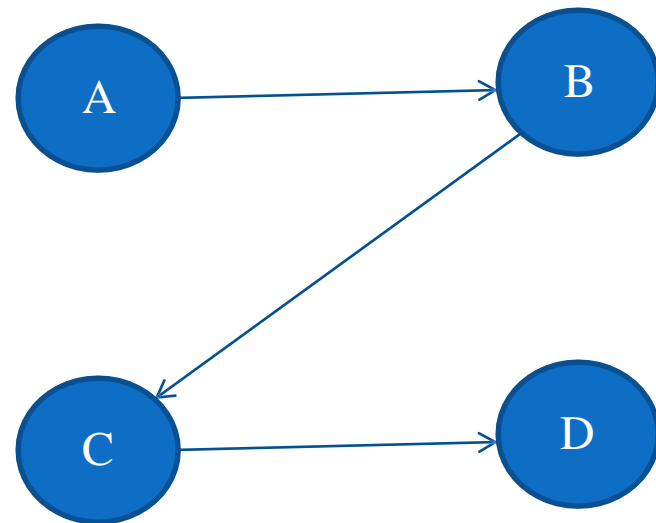
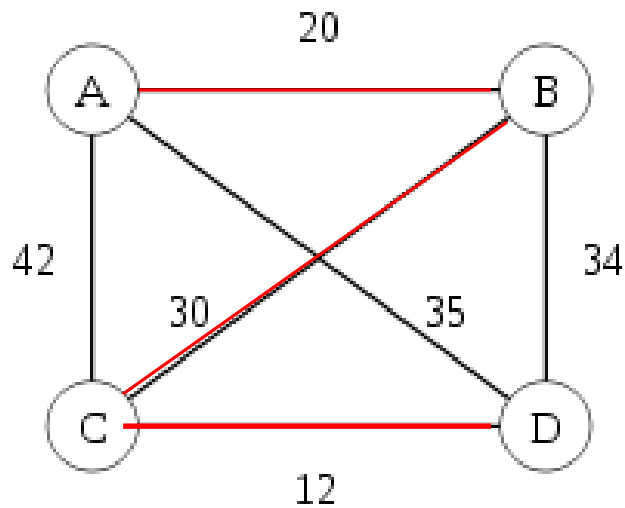


- > Because the nearest city is B, so we go to B

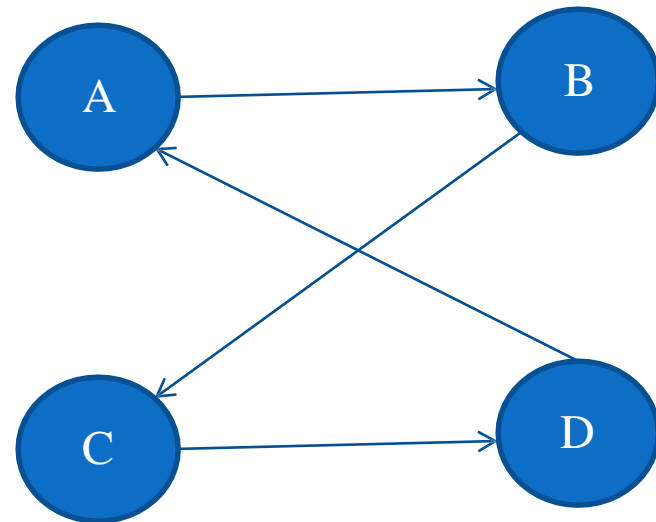
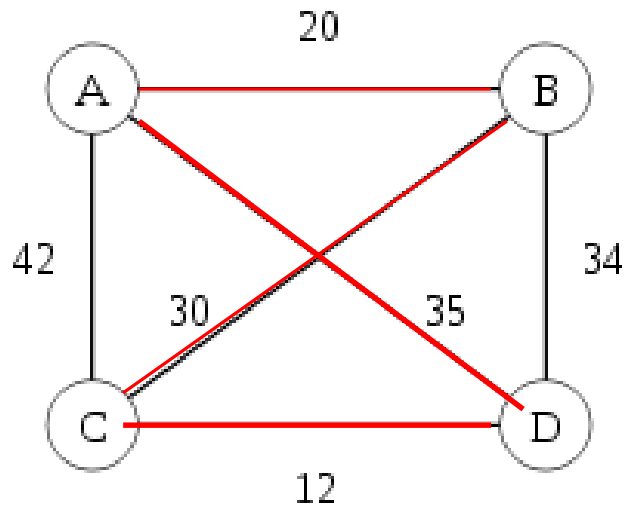
- 
- > From B, we find any other city but A (because A has been visited) that has nearest path. So we choose C:



- 
- > From C, we look to nearest city again, but don't look for A and B, because both have been visited. So we choose D.



> At this node(D), we can't go to any city, because all neighbor of D has been visited. We go back to first city(A).



> And that was how to solve TSP problem.

Advantage of Greedy

- > Greedy is easy to be implemented. Just search the best choice from the current state that 'reachable' (has any paths or any connections).
- > In simple case, greedy often give you the best solution.





Drawback of Greedy

- > In large and complex case, greedy doesn't always give you the best solution, because it's just search and take the best choice that you can reach from the current state.
- > It takes longer time than any other algorithms for big case of problem



Analysis

- A greedy algorithm typically makes (approximately) n choices for a problem of size n
 - (The first or last choice may be forced)
- Hence the expected running time is:
 $O(n * O(\text{choice}(n)))$, where $\text{choice}(n)$ is making a choice among n objects
 - Counting: Must find largest useable coin from among k sizes of coin (k is a constant), an $O(k)=O(1)$ operation;
 - Therefore, coin counting is (n)
 - Huffman: Must sort n values before making n choices
 - Therefore, Huffman is $O(n \log n) + O(n) = O(n \log n)$
 - Minimum spanning tree: At each new node, must include new edges and keep them sorted, which is $O(n \log n)$ overall
 - Therefore, MST is $O(n \log n) + O(n) = O(n \log n)$



Other greedy algorithms

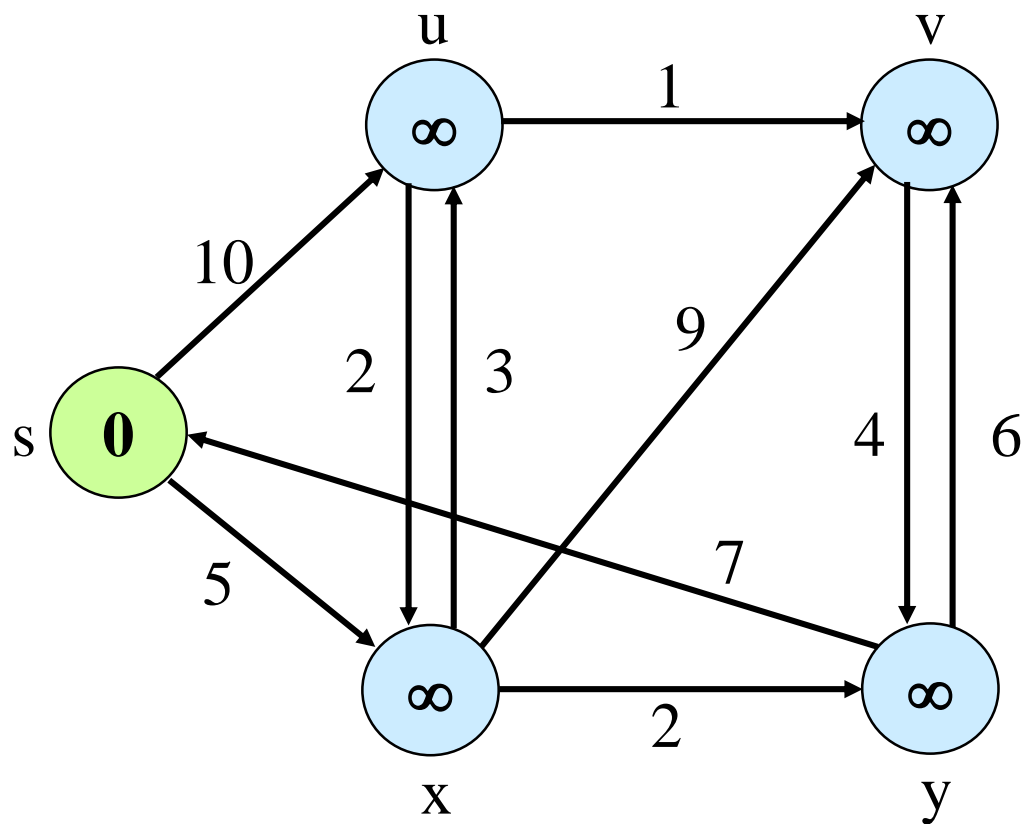
- Dijkstra's algorithm for finding the shortest path in a graph
 - Always takes the *shortest* edge connecting a known node to an unknown node
- Kruskal's algorithm for finding a minimum-cost spanning tree
 - Always tries the *lowest-cost* remaining edge
- Prim's algorithm for finding a minimum-cost spanning tree
 - Always takes the *lowest-cost* edge between nodes in the spanning tree and nodes not yet in the spanning tree



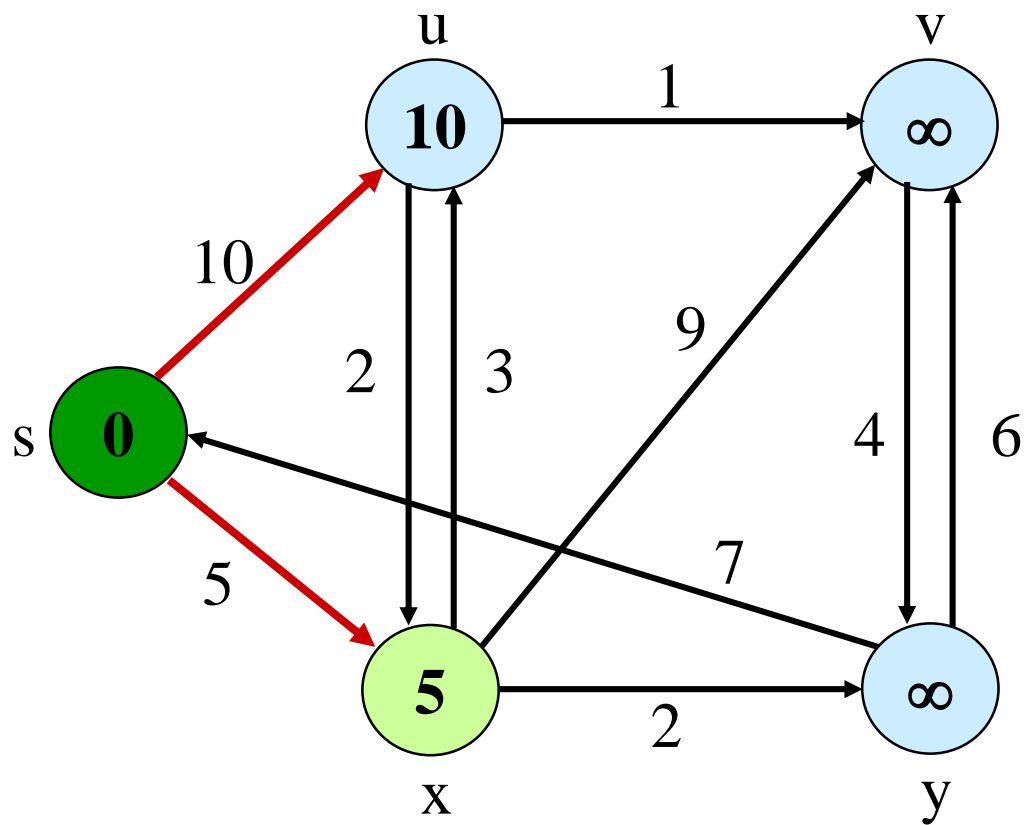
Dijkstra's shortest-path algorithm

- Dijkstra's algorithm finds the shortest paths from a given node to all other nodes in a graph
 - Initially,
 - Mark the given node as *known* (path length is zero)
 - For each out-edge, set the distance in each neighboring node equal to the *cost* (length) of the out-edge, and set its *predecessor* to the initially given node
 - Repeatedly (until all nodes are known),
 - Find an unknown node containing the smallest distance
 - Mark the new node as known
 - For each node adjacent to the new node, examine its neighbors to see whether their estimated distance can be reduced (distance to known node plus cost of out-edge)
 - If so, also reset the predecessor of the new node

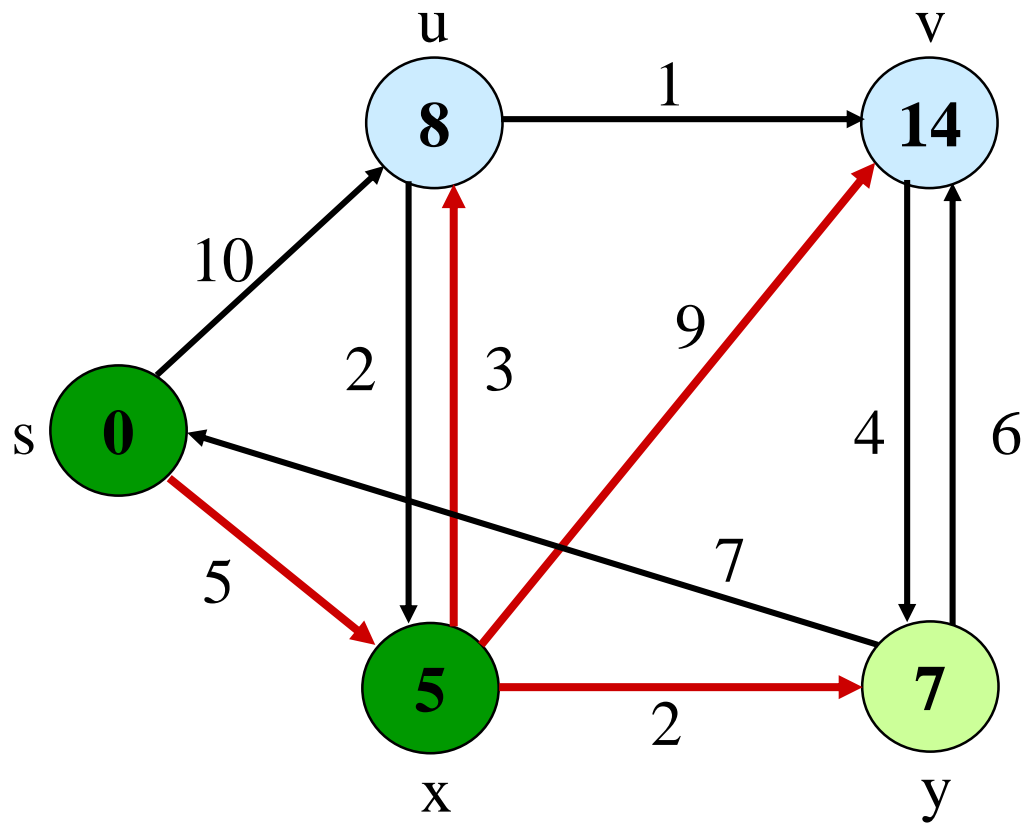
Example



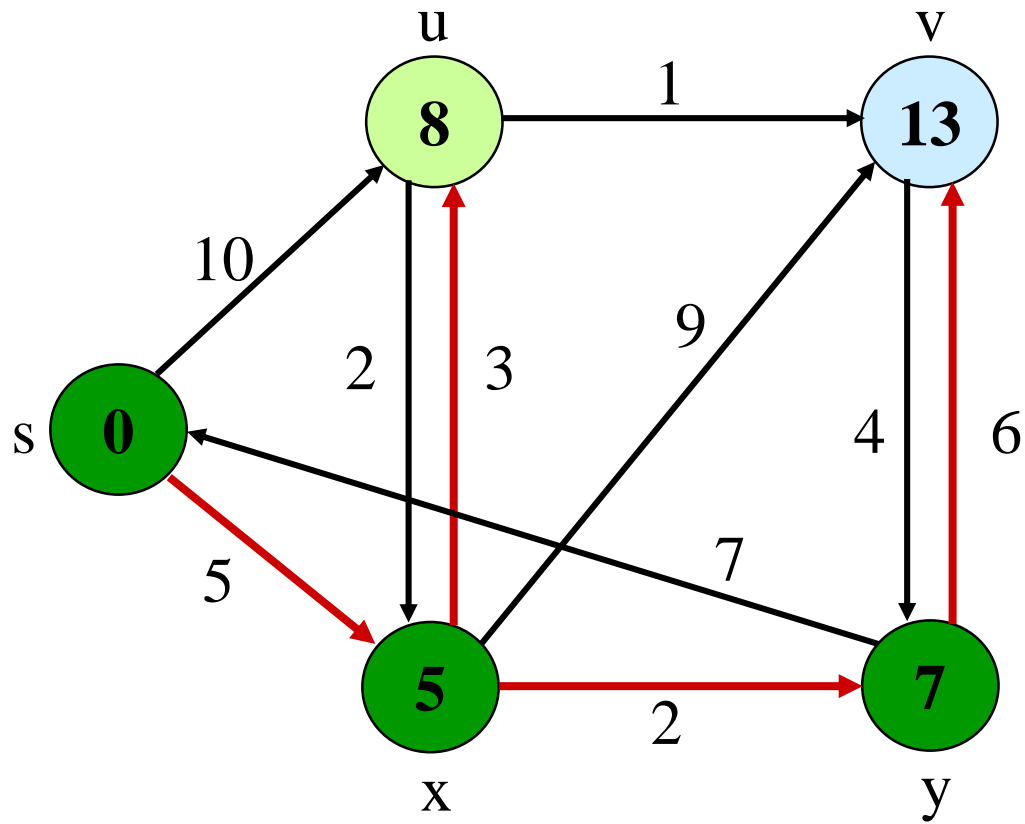
Example



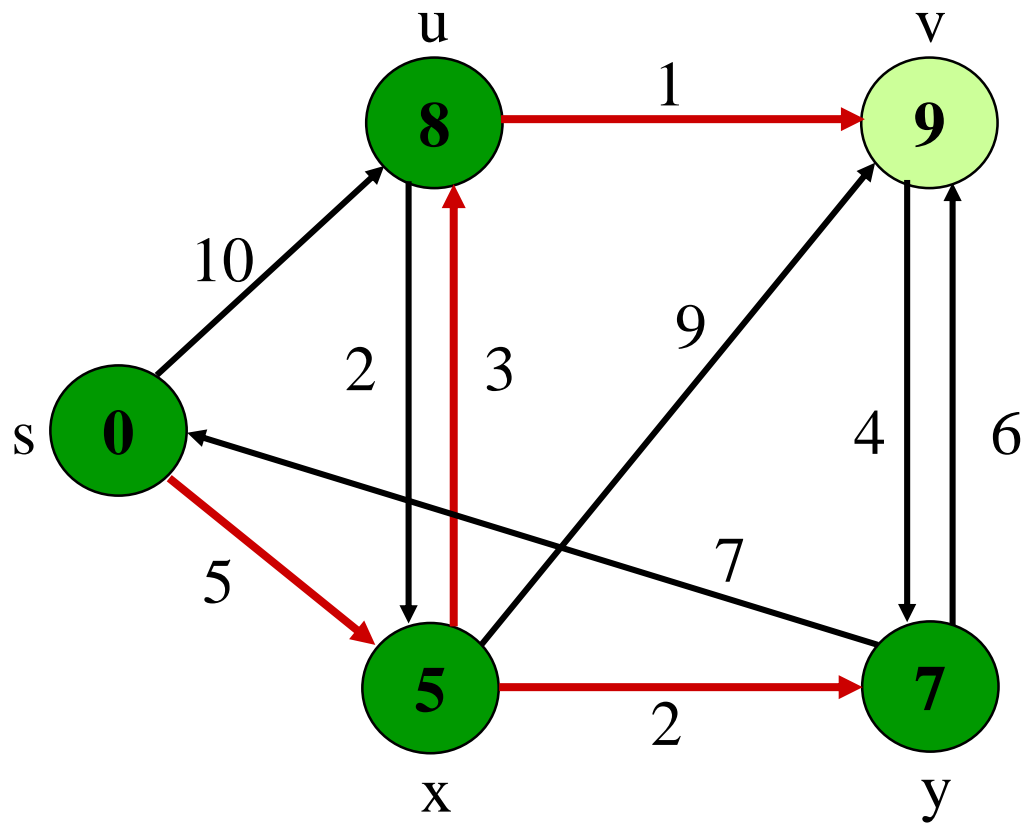
Example



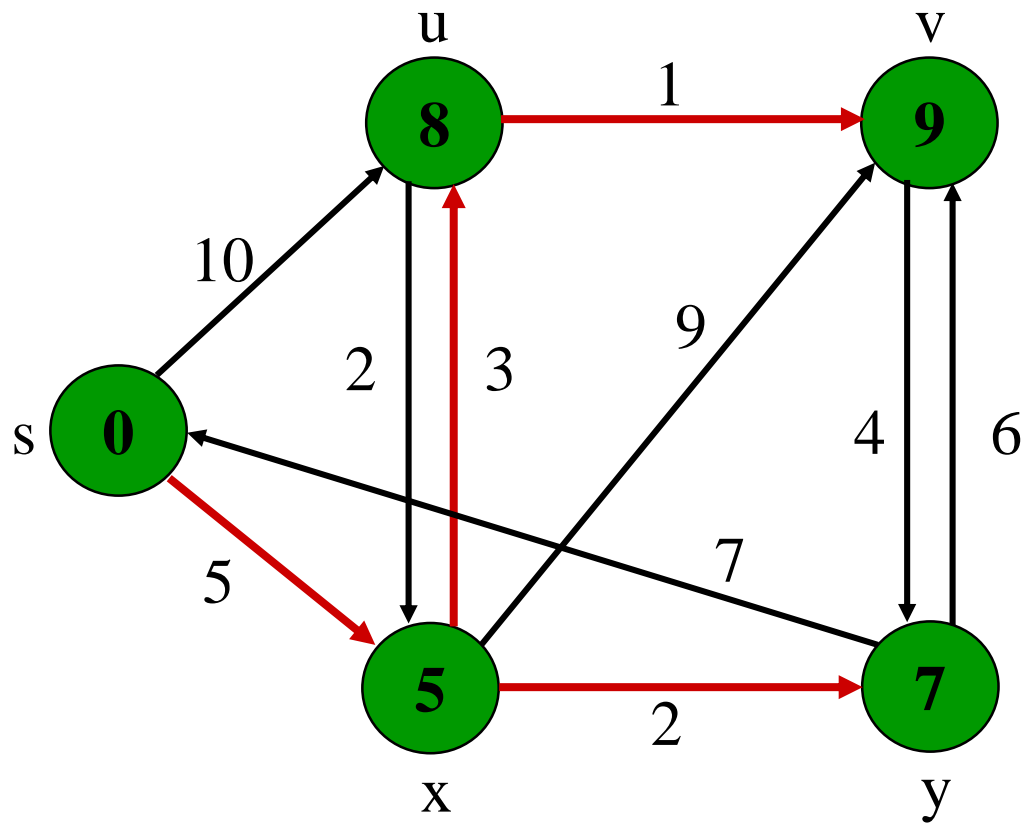
Example



Example



Example





Analysis of Dijkstra's algorithm I

- Assume that the *average* out-degree of a node is some constant k
 - Initially,
 - Mark the given node as *known* (path length is zero)
 - This takes $O(1)$ (constant) time
 - For each out-edge, set the distance in each neighboring node equal to the *cost* (length) of the out-edge, and set its *predecessor* to the initially given node
 - If each node refers to a list of k adjacent node/edge pairs, this takes $O(k) = O(1)$ time, that is, constant time
 - Notice that this operation takes *longer* if we have to extract a list of names from a hash table



Analysis of Dijkstra's algorithm II

- Repeatedly (until all nodes are known), (n times)
 - Find an unknown node containing the smallest distance
 - Probably the best way to do this is to put the unknown nodes into a priority queue; this takes $k * O(\log n)$ time *each* time a new node is marked “known” (and this happens n times)
 - Mark the new node as known -- $O(1)$ time
 - For each node adjacent to the new node, examine its neighbors to see whether their estimated distance can be reduced (distance to known node plus cost of out-edge)
 - If so, also reset the predecessor of the new node
 - There are k adjacent nodes (on average), operation requires constant time at each, therefore $O(k)$ (constant) time
 - Combining all the parts, we get:
 $O(1) + n * (k * O(\log n) + O(k))$, that is, $O(nk \log n)$ time



Kruskal's MST Algorithm

Prepared by
Kenrick Mock

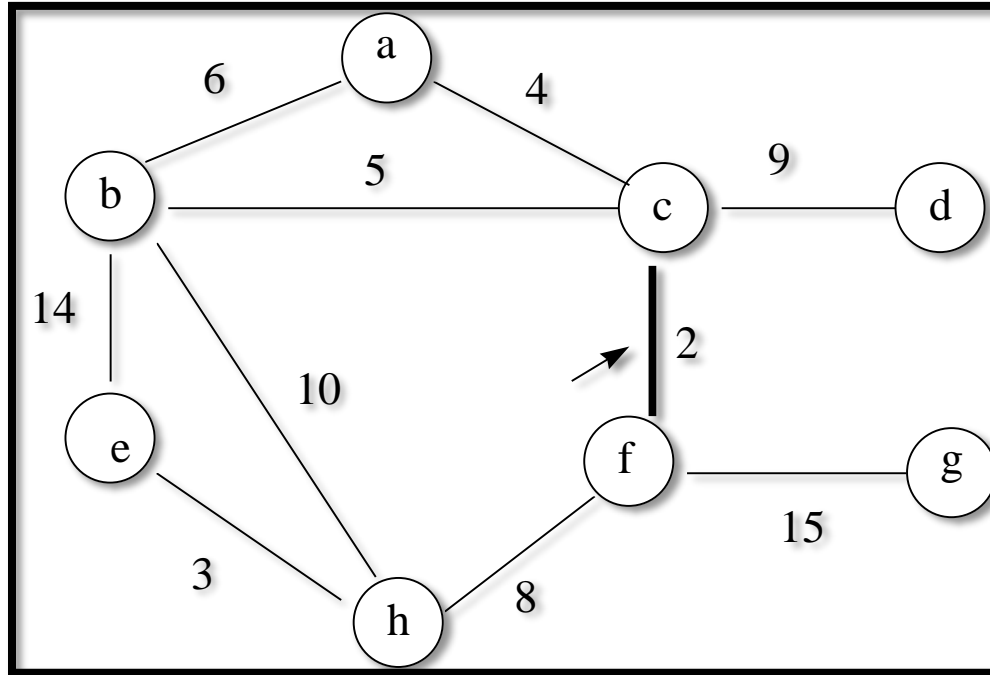
- Idea: Greedily construct the MST
 - Go through the list of edges and make a forest that is a MST
 - At each vertex, sort the edges
 - Edges with smallest weights examined and possibly added to MST before edges with higher weights
 - Edges added must be “safe edges” that do not ruin the tree property.



Kruskal's Algorithm

```
Kruskal( $G, w$ )           ; Graph  $G$ , with weights  $w$ 
   $A \leftarrow \{\}$        ; Our MST starts empty
  for each vertex  $v \in V[G]$  do Make-Set( $v$ )  ; Make each vertex a set
  Sort edges of  $E$  by increasing weight
  for each edge  $(u, v) \in E$  in order
    ; Find-Set returns a representative (first vertex) in the set
    do if Find-Set( $u$ )  $\neq$  Find-Set( $v$ )
      then  $A \leftarrow A \cup \{(u, v)\}$ 
      Union( $u, v$ )           ; Combines two trees
  return  $A$ 
```


Kruskal's Example



- $A = \{ \}$, Make each element its own set. $\{a\} \{b\} \{c\} \{d\} \{e\} \{f\} \{g\} \{h\}$
- Sort edges.
- Look at smallest edge first: $\{c\}$ and $\{f\}$ not in same set, add it to A , union together.
- Now get $\{a\} \{b\} \{c\ f\} \{d\} \{e\} \{g\} \{h\}$

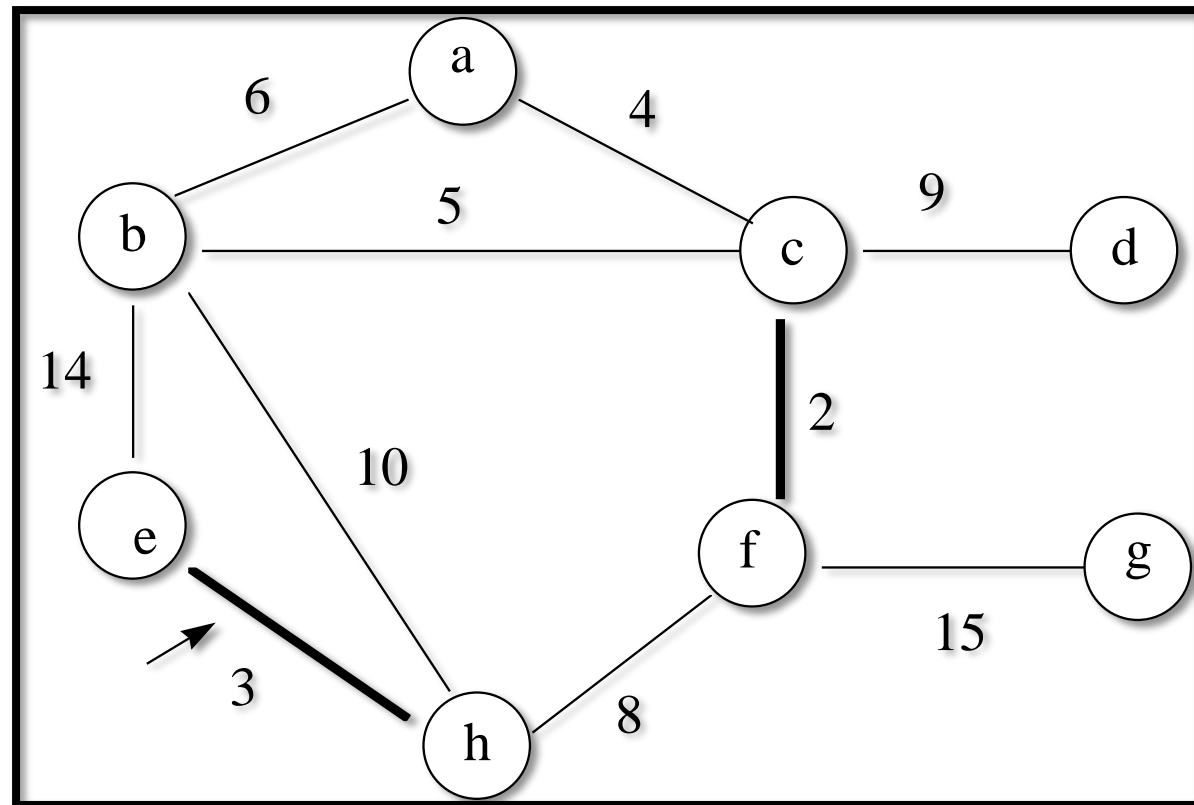
Kruskal's Example

Keep going, checking next smallest edge.

Had: {a} {b} {c f} {d} {e} {g} {h}

{e} \neq {h}, add edge.

Now get {a} {b} {c f} {d}
{e h} {g}



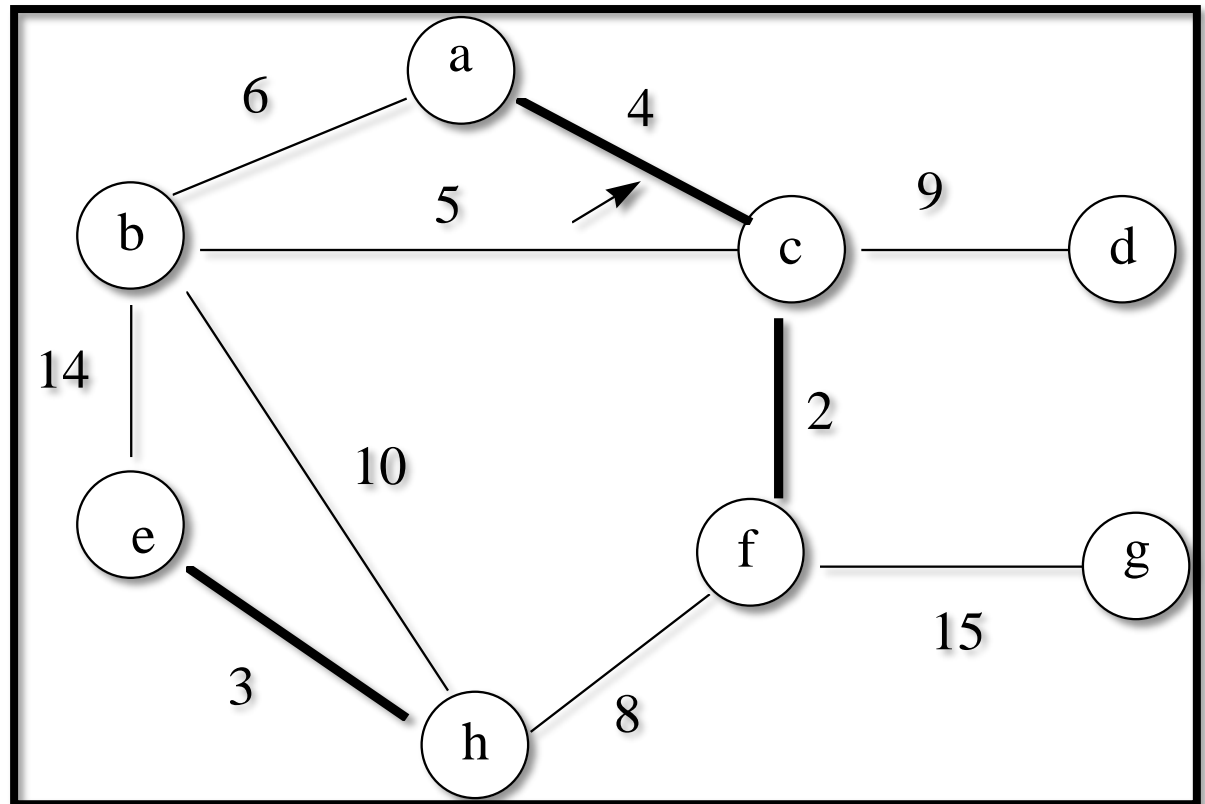
Kruskal's Example

Keep going, checking next smallest edge.

Had: {a} {b} {c f} {d} {e h} {g}

{a} \neq {c f}, add edge.

Now get {b} {a c f}
{d} {e h} {g}

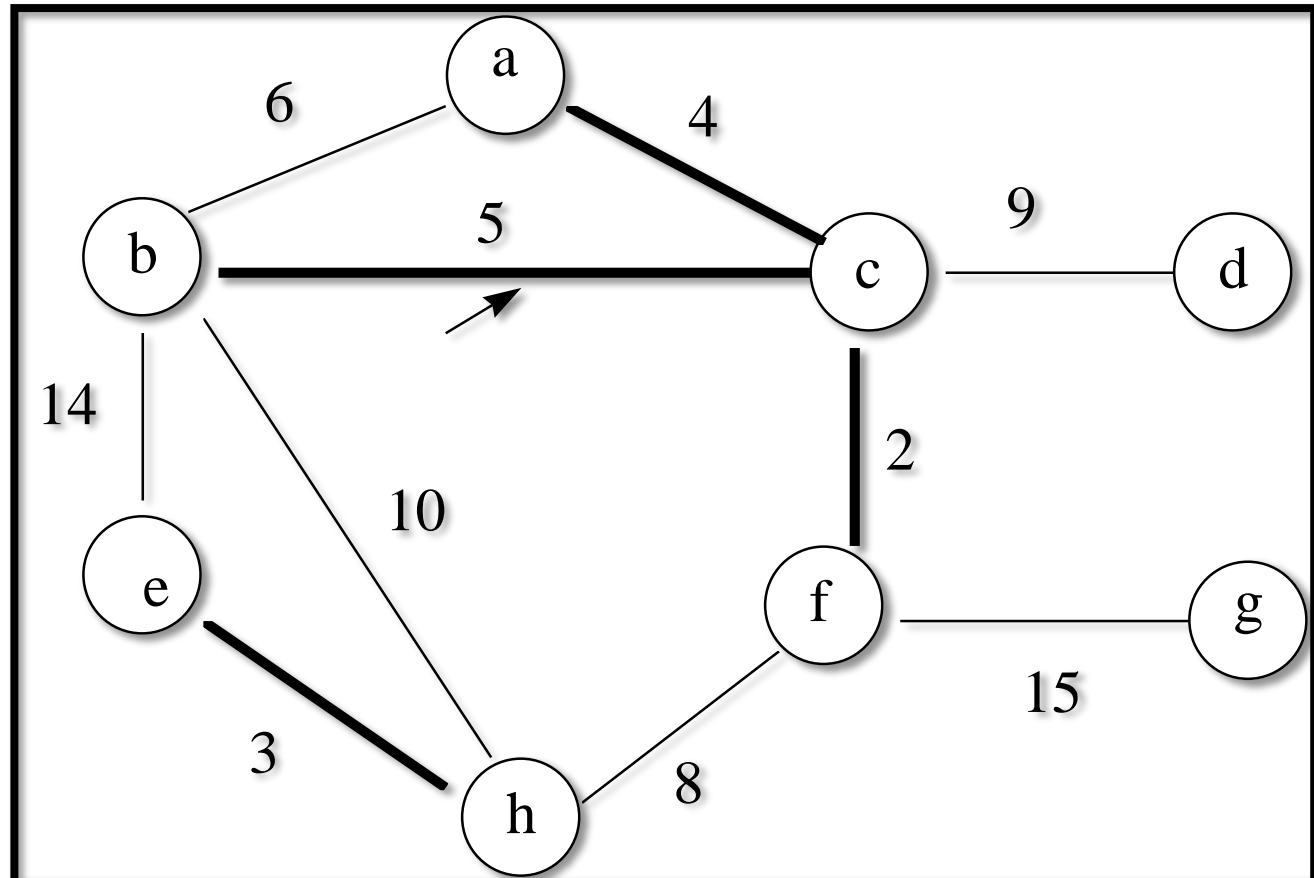


Kruskal's Example

Keep going, checking next smallest edge.

Had $\{b\}$ $\{a\ c\ f\}$ $\{d\}$ $\{e\ h\}$ $\{g\}$

$\{b\} \neq \{a\ c\ f\}$, add edge.



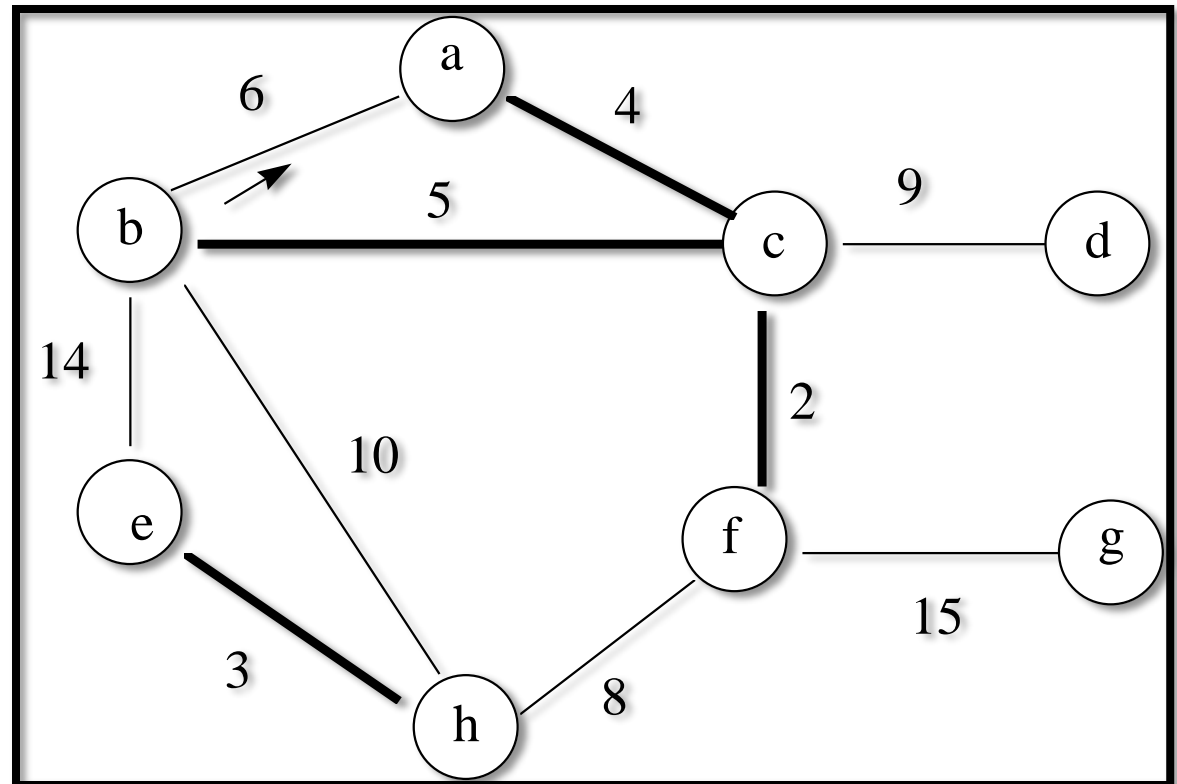
Now get $\{a\ b\ c\ f\}$
 $\{d\}$ $\{e\ h\}$ $\{g\}$

Kruskal's Example

Keep going, checking next smallest edge.

Had $\{a\ b\ c\ f\}$ $\{d\}$ $\{e\ h\}$ $\{g\}$

$\{a\ b\ c\ f\} = \{a\ b\ c\ f\}$, don't add it!



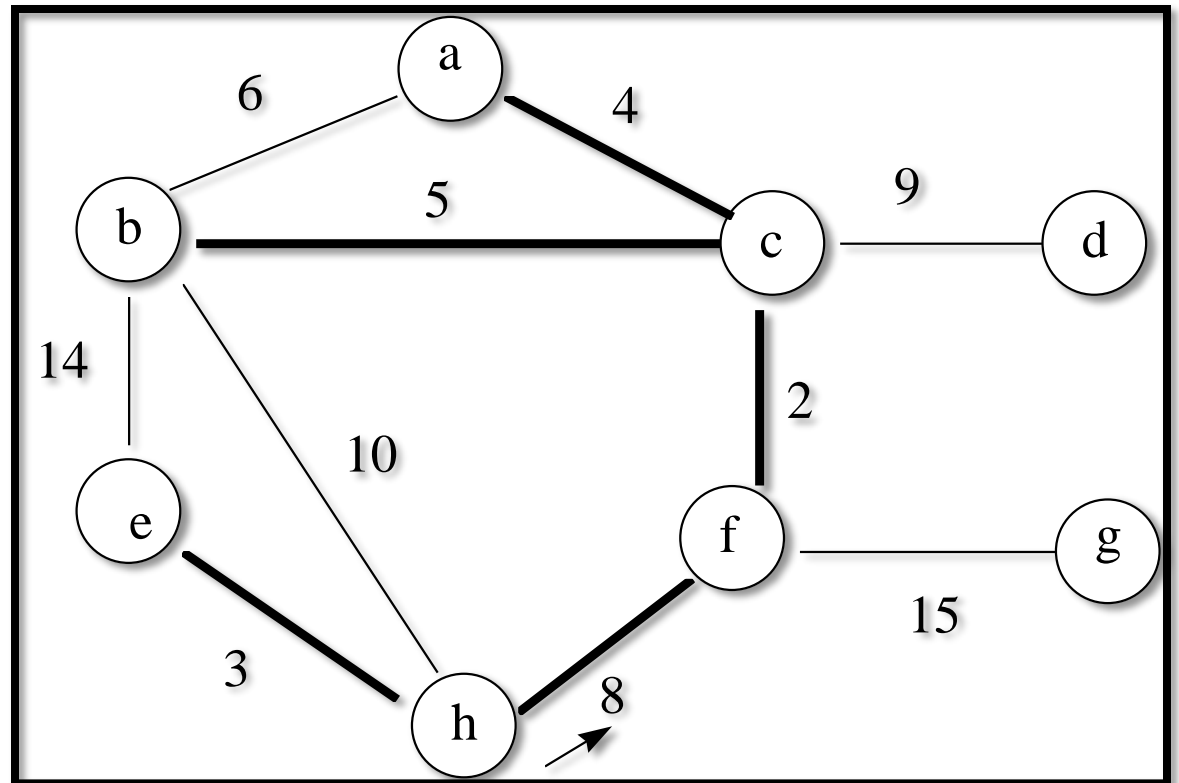
Kruskal's Example

Keep going, checking next smallest edge.

Had $\{a\ b\ c\ f\}$ $\{d\}$ $\{e\ h\}$ $\{g\}$

$\{a\ b\ c\ f\} = \{e\ h\}$, add it.

Now get
 $\{a\ b\ c\ f\ e\ h\}$
 $\{d\}\{g\}$



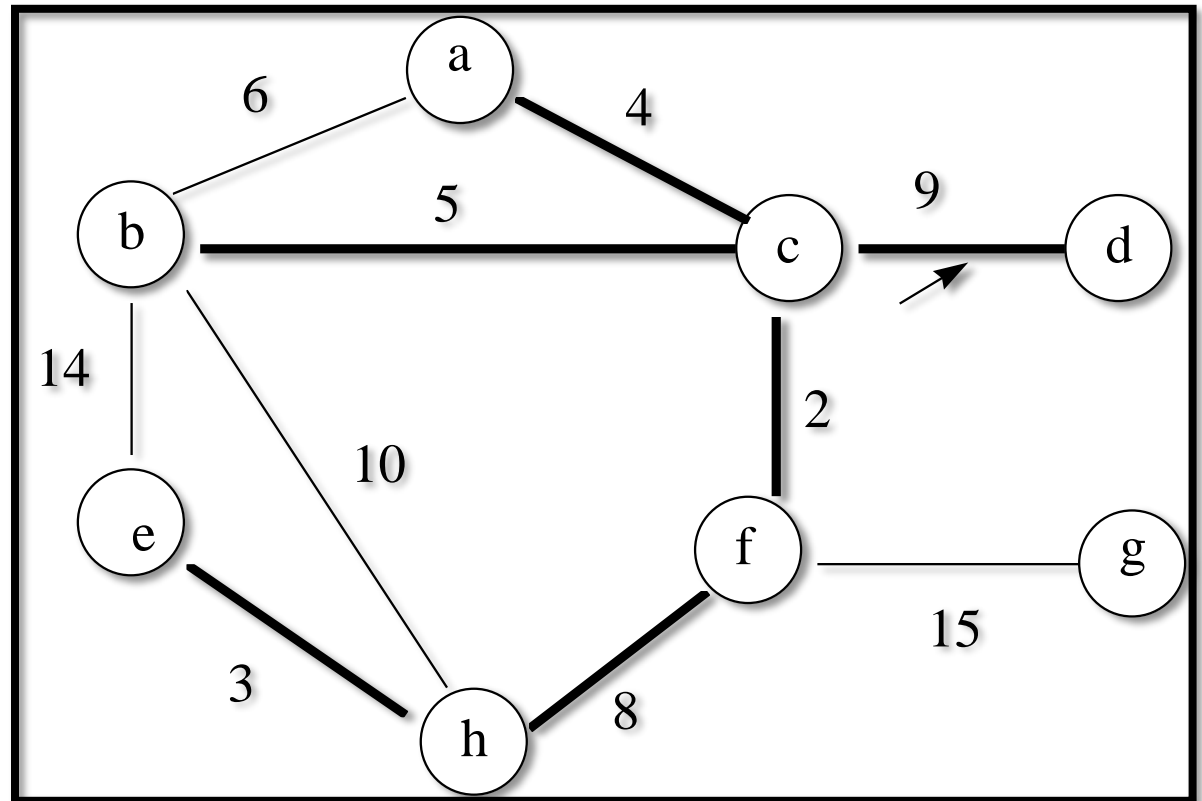
Kruskal's Example

Keep going, checking next smallest edge.

Had $\{a\ b\ c\ f\ e\ h\}\ \{d\}\{g\}$

$\{d\} \neq \{a\ b\ c\ e\ f\ h\}$, add it.

Now get
 $\{a\ b\ c\ d\ e\ f\ h\}$
 $\{g\}$

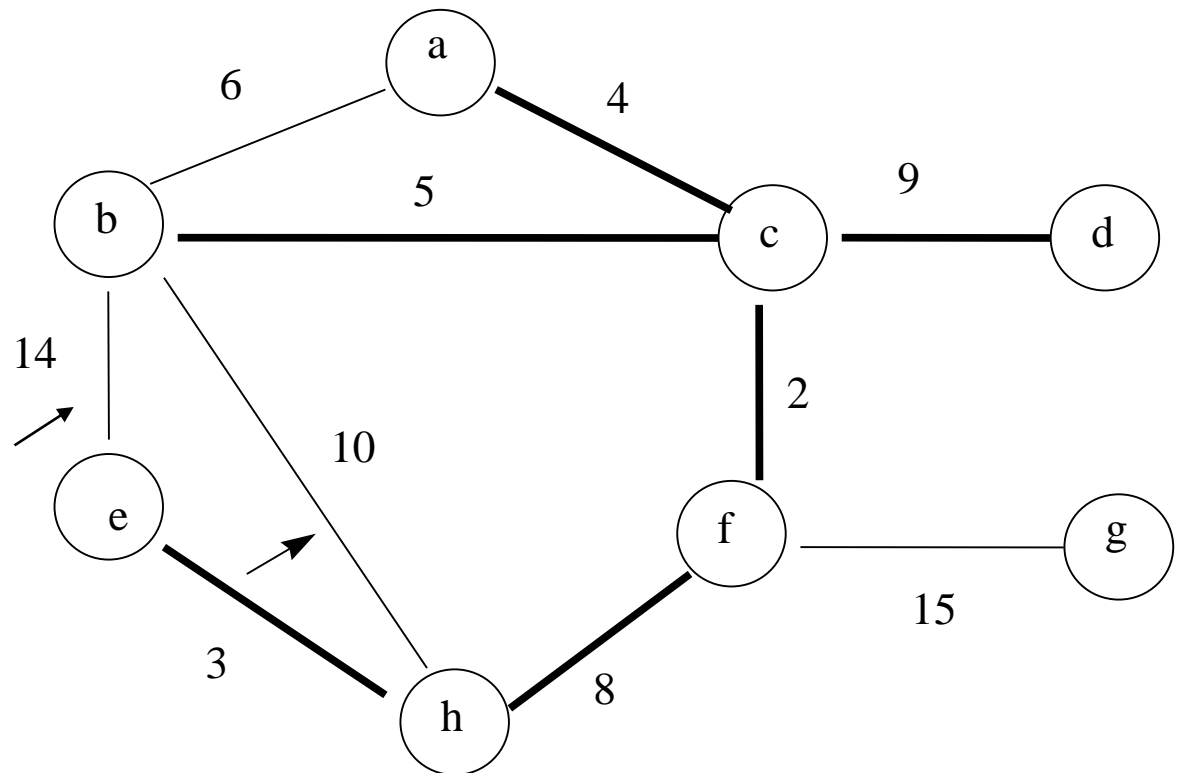


Kruskal's Example

Keep going, check next two smallest edges.

Had $\{a\ b\ c\ d\ e\ f\ h\}\ \{g\}$

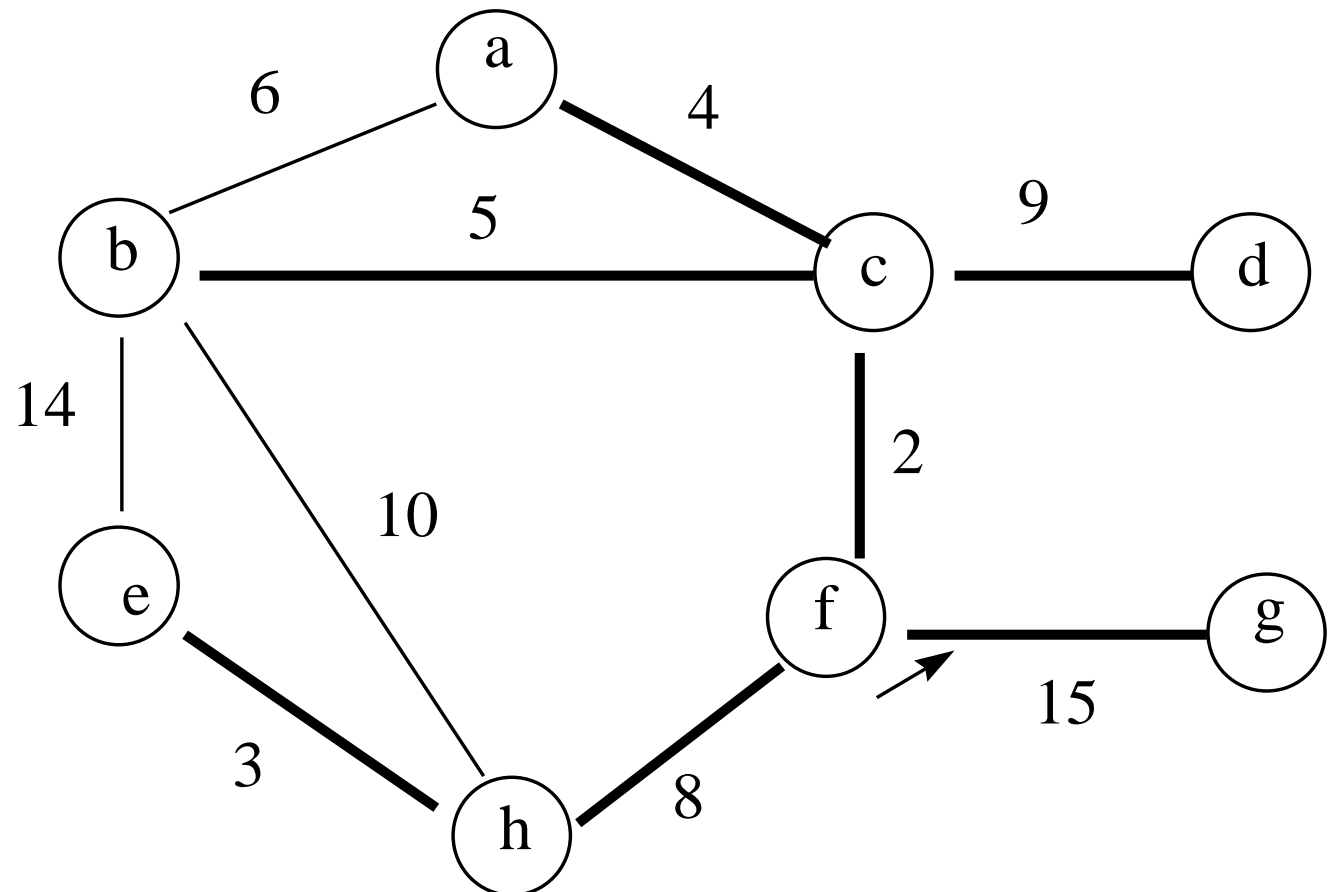
$\{a\ b\ c\ d\ e\ f\ h\} = \{a\ b\ c\ d\ e\ f\ h\}$, don't add it.



Kruskal's Example

Do add the last one:

Had $\{a\ b\ c\ d\ e\ f\ h\}\ \{g\}$





Runtime of Kruskal's Algo

- Runtime depends upon time to union set, find set, make set
- Simple set implementation: number each vertex and use an array
 - Use an array

$\text{member}[]$: $\text{member}[i]$ is a number j such that the i th vertex is a member of the j th set.
 - Example

$\text{member}[1,4,1,2,2]$
indicates the sets $S1=\{1,3\}$, $S2=\{4,5\}$ and $S4=\{2\}$;
i.e. position in the array gives the set number. Idea similar to counting sort, up to number of edge members.

Set Operations

- Given the Member array

- Make-Set(v)

$\text{member}[v] = v$

Make-Set runs in constant running time for a single set.

- Find-Set(v)

Return $\text{member}[v]$

Find-Set runs in constant time.

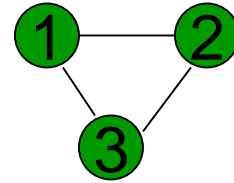
- Union(u, v)

for $i=1$ to n

do if $\text{member}[i] = u$ then $\text{member}[i]=v$

Scan through the member array and update old members to be the new set.

Running time $O(n)$, length of member array.



$\text{member} =$
 $[1, 2, 3] ; \{1\} \{2\} \{3\}$

$\text{find-set}(2) = 2$

Union(2,3)
 $\text{member} = [1, 3, 3] ; \{1\}$
 $\{2\} \{3\}$



Overall Runtime

Kruskal(G, w) ; Graph G , with weights w $O(V)$
 $A \leftarrow \{\}$; Our MST starts empty
 for each vertex $v \in V[G]$ do Make-Set(v) ; Make each vertex a set
 Sort edges of E by increasing weight $O(E \lg E)$ – using heapsort
 for each edge $(u, v) \in E$ in order $O(E)$
 ; Find-Set returns a representative (first vertex) in the set
 do if Find-Set(u) \neq Find-Set(v) $O(1)$
 then $A \leftarrow A \cup \{(u, v)\}$
 Union(u, v) ; Combines two trees $O(V)$
 return A

Total runtime: $O(V) + O(E \lg E) + O(E * (1 + V)) = O(E * V)$

Book describes a version using disjoint sets that runs in $O(E * \lg E)$ time